

Contents

1	Introduction	1
1.1	Motivation	1
1.2	The Virtual Orbital Space Problem	2
1.3	Quantum Computing Context	2
1.4	Objectives	3
1.5	Thesis Outline	3
1.6	Notation and Integral Conventions	3
2	Theory	5
2.1	The Electronic Structure Problem	5
2.1.1	Molecular Orbital Integrals	6
2.1.2	Orbital Space	7
2.1.3	Spin-Orbital Formalism	7
2.2	Wavefunction Methods	8
2.2.1	Hartree-Fock Theory	8
2.2.2	Møller-Plesset Perturbation Theory	9
2.3	Background for understanding OVOS	11
2.3.1	Orbital Rotations and Stationarity	11
2.3.2	The Second-Order Hylleraas Functional	13
2.3.3	Orbital Rotation Parameterisation	16
2.3.4	Gradient Expression	17
2.3.5	Hessian Expression	19
2.3.6	Newton-Raphson	22
2.3.7	Orbital Canonicalisation	22
2.3.8	Computational Cost and Scaling	23
2.4	Relevance to Quantum Computing	25
2.4.1	Brueckner Orbitals and Connection to OVOS	25
2.4.2	Quantum Computing and Orbital Space	26
3	OVOS Implementation	29
3.1	Algorithm	29
3.1.1	Integral Transformation	30
3.1.2	Initial Active Space Selection	30
3.1.3	MP1 Amplitudes and MP2 Energy	30
3.1.4	Gradient and Hessian Assembly	30
3.1.5	Orbital Rotation and Update	30
3.1.6	Convergence and Iteration Control	30
3.2	Initialisation Strategies	30
3.2.1	Default SCF Initialisation	30
3.2.2	Warm-Start from Previous Orbital Set	30
3.2.3	Random Unitary Sampling	30

4	Computational Detail	31
4.0.1	Software and Libraries	31
4.0.2	Molecular Geometries	31
4.0.3	Basis Sets	32
4.1	OVOS Calculation Settings	32
4.2	VQE Calculation Settings	33
5	Results	35
5.1	OVOS	35
5.1.1	Energy Recovery as a Function of Active Virtual Orbital Count	35
5.1.2	Basis Set Dependence	44
5.2	VQE	49
5.2.1	OVOS Orbitals as a VQE Reference State	49
5.2.2	T1 Amplitudes and the Brueckner Orbital Connection	49
5.3	ooVQE	50
6	Conclusion and Outlook	51
6.1	Conclusion	51
6.1.1	Summary of Findings	51
6.1.2	Critical Assessment	51
6.1.3	Implications for Quantum Computing	51
6.2	Outlook	51
6.2.1	Algorithmic Improvements	51
6.2.2	Method Extensions	51
6.2.3	Quantum Computing Applications	51
6.2.4	Larger Molecular Systems	52
6.2.5	Benchmarking and Validation	52
	Bibliography	53
A	Implementation Validation	55
B	Additional Data and Figures	57
C	Implementation Details and Source Code	59

Chapter 1

Introduction

This chapter introduces the motivation, context, and objectives of this thesis, which focuses on the optimisation of virtual orbital spaces for post-Hartree-Fock methods and their implications for quantum computing. We begin by outlining the computational challenges in quantum chemistry that motivate the need for orbital optimisation. We then review existing strategies for virtual orbital reduction, leading to the OVOS method as the focus of this work. Finally, we state the specific research objectives and provide an outline of the thesis structure.

"NEED TO CHECK WITH PROJECT DESCRIPTION"

1.1 Motivation

Chemistry fundamentally arises from the quantum mechanical interactions of electrons and nuclei. The exact solution of the electronic Schrödinger equation for a molecule with N electrons scales as $\mathcal{O}(N!)$ in the size of the configurational space, which rapidly renders exact calculations (such as Full Configuration Interaction) infeasible for any but the smallest systems [1]. To achieve chemical accuracy (approximately 1 kcal/mol), it is essential to account for electron correlation effects beyond the mean-field Hartree-Fock (HF) approximation. Post-HF methods such as second-order Møller-Plesset perturbation theory (MP2), coupled cluster with single and double excitations (CCSD), and CCSD with perturbative triples [CCSD(T)] are widely used to capture these correlation effects [2]. However, the computational cost of these methods scales steeply with system size, primarily due to the virtual orbital space. The number of virtual orbitals grows with the size of the basis set, and the computational cost of post-HF methods scales as $\mathcal{O}(O^2V^2)$ for MP2 (AO to MO basis), $\mathcal{O}(V^4O^2)$ for CCSD, and $\mathcal{O}(O^3V^4)$ for CCSD(T), where O is the number of occupied and V is the number of virtual orbitals, notation from [3]. As occupied orbitals are typically fewer than virtual orbitals, the virtual space dominates the computational cost.

This challenge is not only relevant for classical computational chemistry but also has direct implications for quantum computing. Quantum algorithms for chemistry, such as the Variational Quantum Eigensolver (VQE)[4], require a mapping of the electronic structure problem onto qubits [4]. The number of qubits required scales with the number of orbitals. Although current NISQ devices may have over 100 physical qubits (e.g., IBM's systems with 130+ qubits), device noise severely limits the number of qubits that can be reliably used. Reducing the number of orbitals (the virtual orbital space) is therefore crucial to minimise qubit requirements and mitigate the effects of noise, enabling chemistry simulations feasible in the near term [5]. This motivates a closer examination of the virtual orbital space itself, which is the bottleneck for all post-HF methods and is the subject of the following section.

1.2 The Virtual Orbital Space Problem

The virtual orbital space, which consists of the unoccupied molecular orbitals, plays a critical role in post-Hartree-Fock methods as it defines the space of excited configurations that contribute to electron correlation. The number of virtual orbitals grows with the size of the basis set, and their contribution to the correlation energy is not uniform. This section surveys various strategies that have been developed to reduce or optimise the virtual orbital space, ultimately motivating the focus of this thesis on the Optimised Virtual Orbital Space (OVOS) method, originally proposed by Adamowicz and Bartlett in 1987 [6].

The standard approach to defining virtual orbitals is to use the canonical Hartree-Fock (HF) virtual orbitals, which are obtained by diagonalising the Fock matrix [3]. These orbitals are optimal for describing the single-determinant HF state. However, when we extend to multi-determinantal post-HF methods such as configuration interaction (CI) or coupled cluster (CC), the HF orbitals are no longer optimal for capturing electron correlation. This motivates the need to optimize the orbital space for correlated wavefunctions.

Natural orbitals, introduced by Löwdin in 1955 [7], are obtained by diagonalising the one-particle reduced density matrix (1-RDM) of an MP2 wavefunction. They are ordered by their occupation numbers, which can be used to truncate the virtual space. However, this approach requires an MP2 calculation to obtain the 1-RDM, creating a circular dependency.

Frozen natural orbitals (FNO) are a practical solution to this problem. They are computed from a low-cost correlated calculation, such as MP2, and the virtual orbitals are truncated based on an occupation threshold [8]. This method is widely used in practice and can significantly reduce the computational cost while retaining accuracy.

The OVOS method takes a different approach by directly optimising a subset of the virtual orbitals to maximise the second-order correlation energy [6]. Unlike natural orbitals and FNO, OVOS does not require an initial MP2 calculation and can be systematically applied to any desired number of virtual orbitals. The OVOS method is the focus of this thesis, and we will implement and benchmark it in the context of unrestricted Hartree-Fock (UHF) calculations using the PySCF software package [9].

1.3 Quantum Computing Context

Quantum computing has emerged as a promising paradigm for simulating quantum systems, including molecular electronic structure. Algorithms such as VQE can efficiently find ground state energies by preparing a parameterised quantum state and optimising the parameters using a classical optimiser [4]. However, current NISQ devices are severely limited in the number of qubits and the depth of quantum circuits they can reliably execute, primarily due to noise [5]. This makes the connection between orbital-space reduction and quantum resource requirements directly relevant to this thesis.

In quantum chemistry, the electronic structure problem is typically mapped onto qubits using transformations such as the Jordan-Wigner [10], which requires $Q = 2N_{\text{orb}}$ qubits for N_{orb} spatial orbitals. Given the limited number of qubits and the need

for shallow circuits, reducing the number of virtual orbitals directly translates to a reduction in qubit requirements and circuit depth. The OVOS method, by optimising a subset of the virtual orbitals, can significantly reduce the number of qubits needed for a given level of accuracy, $Q = 2N'_{\text{orb}}$, where $N'_{\text{orb}} < N_{\text{orb}}$ is the number of orbitals retained after OVOS optimisation.

1.4 Objectives

The thesis will seek to achieve the following specific objectives:

1. Implement and validate the OVOS algorithm of Adamowicz & Bartlett [6] using the PySCF software package, for a representative set of molecules and basis sets.
2. Characterise the energy recovery as a function of the active virtual orbital count for all test systems, quantifying the fraction of correlation energy recovered by the OVOS-optimised orbitals.
3. Investigate whether OVOS-optimised molecular orbitals satisfy the Brueckner orbital condition at the MP2 level (vanishing T1 amplitudes), and assess the theoretical connection between OVOS and orbital-optimised MP2 [11].
4. Quantify the potential benefit of OVOS-optimised orbitals as a reference trial state for VQE-type quantum computing algorithms: assess the expected improvement in reference-state overlap and compare initialisation strategies (default UHF, OVOS, UMP2 natural orbitals).

1.5 Thesis Outline

"..."

1.6 Notation and Integral Conventions

We establish the following specific notation and conventions for OVOS, maintaining consistency with the original Adamowicz & Bartlett paper [6] .

Index Ordering and Orbital Partitioning

Orbital indices follows the standard convention:

- $i, j, k, l, \dots$ for occupied orbitals
- $a, b, c, d, \dots$ for unoccupied orbitals
- $p, q, r, s, \dots$ for unrestricted orbitals

In OVOS, the virtual orbitals are further partitioned into active ($a, b, \dots$) and inactive ($e, f, \dots$) subsets, with the active space containing N' orbitals optimised for correlation recovery.

Integral Notation: Antisymmetrised vs Non-antisymmetrised

For integral notation, we distinguish between the antisymmetrised physicists' notation $\langle ab||ij\rangle$ and the non-antisymmetrised chemists' notation $\langle ab|ij\rangle$. The antisymmetrised integral is defined as:

$$\langle ab||ij\rangle = \langle ab|ij\rangle - \langle ab|ji\rangle$$

where $\langle ab|ij\rangle$ is the standard two-electron integral in the MO basis. The physicists' notation is used throughout the OVOS derivations for consistency with the original literature and to simplify the expressions for the gradient and Hessian.

Integral Notation: Chemists' vs. Physicists' Conventions

The notation for two-electron integrals can differ between chemists and physicists, leading to potential confusion when comparing formulas from the literature with the output of quantum chemistry software. In chemists' notation, the integral is written as $(pq|rs)$, while in physicists' notation, it is written as $\langle pq|rs\rangle = (pr|qs)$ [3]. The antisymmetrised integral, which is relevant for post-HF methods, is defined as:

$$\langle pq||rs\rangle = \langle pq|rs\rangle - \langle pq|sr\rangle = (pr|qs) - (ps|qr)$$

The OVOS formulas in this thesis use the antisymmetrised physicists' notation, while PySCF returns integrals in the chemists' notation[9]. The conversion between these notations is explicitly implemented in the code as: $\langle pq||rs\rangle = (pr|qs) - (ps|qr)$. This establishes a clear and consistent framework for interpreting the integrals used in the OVOS method and ensures that the theoretical formulas align with the computational implementation.

Chapter 2

Theory

This chapter is meant to provide the theoretical background necessary to understand the OVOS method and its implementation, and how we extend it to Quantum computing. The content is structured to build from fundamental concepts in electronic structure theory to the specific mathematical formulations used in OVOS. The chapter is structured in four sections: (2.1) the electronic structure problem and the importance of electron correlation, (2.2) the wavefunction methods that are the basis for OVOS, (2.3) the mathematical framework for OVOS, and (2.4) the relevant aspects of OVOS on quantum computers.

2.1 The Electronic Structure Problem

The electronic structure problem is the task of solving the non-relativistic, time-independent Schrödinger equation for a system of N electrons moving in the electrostatic potential of M fixed nuclei, reflecting the Born-Oppenheimer approximation.¹ Typically, the electronic Hamiltonian is expressed as the Schrödinger equation [3]:

$$\hat{H}_{el}\Psi(\mathbf{r}_1, \mathbf{r}_2, \dots, \mathbf{r}_N) = E\Psi(\mathbf{r}_1, \mathbf{r}_2, \dots, \mathbf{r}_N) \quad (2.1)$$

where $\hat{H}_{el}$ is the electronic Hamiltonian, Ψ is the many-electron wavefunction, and E is the electronic energy. The electronic Hamiltonian can be written as [3]:

$$\hat{H}_{el} = \sum_{i=1}^N \left(-\frac{1}{2} \nabla_i^2 - \sum_{A=1}^M \frac{Z_A}{|\mathbf{r}_i - \mathbf{R}_A|} \right) + \sum_{i < j} \frac{1}{|\mathbf{r}_i - \mathbf{r}_j|} \quad (2.2)$$

where the first term represents the kinetic energy of the electrons, the second term is the nuclear attraction, and the third term is the electron-electron repulsion.

The electron-electron repulsion term couples all electrons together, making the exact solution of the Schrödinger equation computationally intractable for all but the smallest systems. This motivates the development of approximate methods that can capture the essential physics of electron correlation while remaining computationally feasible [12, 13].

Hartree-Fock theory provides a mean-field approximation to the electronic structure problem [12, 13]. The residual correlation energy, defined as $E_{corr} = E_{exact} - E_{HF}$, quantifies the error of the Hartree-Fock approximation, accurately capturing only a fraction of the total correlation energy [3].

For small molecules, the correlation energy can be on the order of tenths of a Hartree, and can be critical for accurately describing chemical phenomena such as bond dissociation and the properties of open-shell molecules [5]. The following section introduces Hartree-Fock theory in more detail, setting the stage for the post-Hartree-Fock methods that are the focus of this thesis.

¹The Born-Oppenheimer approximation assumes that the nuclei are stationary due to their much larger mass compared to electrons, allowing the electronic and nuclear motions to be separated.




Figure 2.1: Potential energy curves for H_2 illustrating the failure of HF near dissociation and the importance of electron correlation. The HF curve deviates significantly from the exact FCI curve as the bond is stretched, while MP2 captures much of the correlation energy and follows the FCI curve more closely around the equilibrium bond length. The plot is drawn using Fig. 5.20 and 5.12 from [3].

2.1.1 Molecular Orbital Integrals

At the core of electronic structure methods are the one- and two-electron integrals evaluated in the molecular orbital basis. These integrals form the building blocks of the Roothaan-Hall equations and the correlation energy expressions used in post-HF methods [3]. This section defines these integrals and establishes the notation conventions applied consistently throughout this thesis, resolving an important ambiguity that arises when comparing PySCF output with the formulas of Adamowicz & Bartlett [6].

One-Electron Integrals

One-electron integrals reflect the kinetic energy of electrons and their attraction to the nuclei [3]. In an atomic orbital (AO) basis, the overlap integral is defined as: $S_{\mu\nu} = \langle \mu | \nu \rangle = \int \phi_\mu^*(\mathbf{r}) \phi_\nu(\mathbf{r}) d\mathbf{r}$, the kinetic energy integral is: $T_{\mu\nu} = \langle \mu | -\frac{1}{2} \nabla^2 | \nu \rangle$, and the nuclear attraction integral is: $V_{\mu\nu} = \langle \mu | -\sum_A \frac{Z_A}{|\mathbf{r} - \mathbf{R}_A|} | \nu \rangle$ [3]. The core Hamiltonian in the AO basis is then given by: $h_{\mu\nu} = T_{\mu\nu} + V_{\mu\nu}$, see Eq. 2.2. To transform these integrals to the molecular orbital (MO) basis, we use the MO coefficients from the HF calculation: $h_{pq} = \sum_{\mu\nu} C_{\mu p}^* h_{\mu\nu} C_{\nu q}$, where $C_{\mu p}, C_{\nu q}$ are elements of the coefficient matrix C , that transforms from the AO to MO basis. This transformation is essential for evaluating the Fock matrix and the correlation energy in the MO basis. The one-electron integrals are straightforward to compute [9] and are typically stored in a dense matrix format.

Two-Electron Integrals

Two-electron integrals, also known as electron repulsion integrals (ERIs) are defined as: $(pq|rs) = \iint \phi_p^*(1) \phi_q(1) \frac{1}{r_{12}} \phi_r^*(2) \phi_s(2) d1 d2$ in chemists' notation, which is the convention used by PySCF [9]. These integrals represent the Coulombic repulsion between pairs of electrons and are central to the construction of the Fock matrix and the evaluation of correlation energies [3]. The ERIs exhibit an 8-fold permutation symmetry for real orbitals, which can be exploited to reduce computational cost. This allows for the storage of only unique integrals, for indexes $p \geq q$ and $r \geq s$, and $(pq|rs) = (qp|rs) = (pq|sr) = (qp|sr) = (rs|pq) = (sr|pq) = (rs|qp) = (sr|qp)$. The Fock matrix in the MO basis can be expressed in terms of these integrals as: $F_{pq} = h_{pq} + \sum_i [2(p_i|q_i) - (pq|ii)]$ for RHF, where the sum runs over occupied orbitals. The two-electron integrals are typically stored in a 4-dimensional array, and their evaluation is the computational bottleneck in many electronic structure methods [13]. For UHF calculations, the Fock matrix has separate contributions from α and β electrons, leading to a larger number of integrals to be evaluated and stored.

2.1.2 Orbital Space

A quantum state in electronic structure theory is described by a wavefunction that depends on the coordinates of all electrons, and can be represented by Slater determinants, which are products of spatial orbitals and spin functions,

$$\chi_p(\mathbf{r}, m_s) = \phi_p(\mathbf{r})\sigma(m_s) \quad (2.3)$$

where $\phi_p(\mathbf{r}) = \sum_{\mu} C_{\mu p} \phi_{\mu}(\mathbf{r})$ is the spatial orbital expanded in an atomic basis [3]. The spatial part of the spin-orbitals come from basis sets, which are typically composed of Gaussian-type orbitals (GTOs) or Slater-type orbitals (STOs) [13]. The choice of basis set can significantly affect the accuracy and computational cost of electronic structure calculations. Commonly used basis sets include minimal basis sets like STO-3G, split-valence basis sets like 6-31G [14–16], and correlation-consistent basis sets like cc-pVDZ [17]. The size and quality of the basis set determine the number of spin-orbitals available for constructing the wavefunction, which in turn affects the description of electron correlation and the convergence of post-HF methods [3, 13].

2.1.3 Spin-Orbital Formalism

The spin-orbital formalism provides a framework for unrestricted calculations by treating each spatial orbital as giving rise to a pair of spin-orbitals, one for each spin projection (α and β) [3]. A spin-orbital is defined as $\chi_{p,\sigma}(\mathbf{r}, m_s) = \phi_p(\mathbf{r})\sigma(m_s)$, where $\phi_p(\mathbf{r})$ is the spatial part of the orbital and $\sigma(m_s)$ is the spin function, which can be either α (spin-up) or β (spin-down). The total number of spin-orbitals is therefore $n_{spin} = 2 \times n_{spatial}$. In this implementation, an interleaved ordering convention is used for the spin-orbitals: even indices correspond to α spin-orbitals, while odd indices correspond to β spin-orbitals, as illustrated in Figure 2.2. The total number of spatial orbitals is $n_{spatial} = n_{spin}/2$, and the mapping is as follows: index 0 corresponds to spatial orbital 0 with α spin, index 1 corresponds to spatial orbital 0 with β spin, index 2 corresponds to spatial orbital 1 with α spin, etc.

This interleaved ordering allows for a unified treatment of α and β electrons in the construction of the Fock matrix and the evaluation of integrals. The Fock matrix in the spin-orbital basis is block-diagonal, with separate blocks for α and β spins. $F_{spin} = \begin{pmatrix} F_\alpha & 0 \\ 0 & F_\beta \end{pmatrix}$ where F_α and F_β are the Fock matrices for the α and β spin-orbitals, respectively [3].

The two-electron integrals in the spin-orbital basis are also constructed according to this interleaved ordering, with specific blocks corresponding to different spin combinations (e.g., $\alpha\alpha$, $\beta\beta$, $\alpha\beta$) [3]. After the construction of the Fock matrix and the integrals in the spin-orbital basis, the spin-orbitals are re-sorted by their orbital energies to form the canonical ordering used in all subsequent summations and calculations [3]. This ensures that the occupied and virtual orbitals are correctly identified according to their energies, which is essential for the accurate evaluation of correlation energies and amplitudes in post-HF methods [13].

Spatial Orbital	α Index	β Index	Interleaved Sequence	Parity
$p = 0$	0	1	$\chi_{0,\alpha}, \chi_{0,\beta}$	even, odd
$p = 1$	2	3	$\chi_{1,\alpha}, \chi_{1,\beta}$	even, odd
$p = 2$	4	5	$\chi_{2,\alpha}, \chi_{2,\beta}$	even, odd
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$p = n - 1$	$2n - 2$	$2n - 1$	$\chi_{n-1,\alpha}, \chi_{n-1,\beta}$	even, odd

Figure 2.2: Interleaved α/β spin-orbital ordering. Even spin-orbital indices correspond to α (spin-up) orbitals, while odd indices correspond to β (spin-down) orbitals. The total number of spin-orbitals is $n_{spin} = 2n_{spatial}$.

2.2 Wavefunction Methods

...

2.2.1 Hartree-Fock Theory

Hartree-Fock (HF) theory approximates the N-electron wavefunction as a single Slater determinant of N orthonormal spin-orbitals, treating electron-electron repulsion in a mean-field sense [3]. Despite its simplicity, HF underpins virtually all post-HF correlated methods.

Slater determinants are described by the Pauli exclusion principle, ensuring that no two electrons can occupy the same quantum state. This antisymmetry is crucial for accurately describing fermionic systems. The variational principle leads to the Fock equations, which can be expressed in matrix form as the Roothaan-Hall equations[3]: $FC = SC\epsilon$, where F is the Fock matrix, S is the overlap matrix, C is the coefficient matrix for the molecular orbitals, and ϵ is the diagonal matrix of orbital energies. The Fock matrix incorporates the Coulomb (J) and exchange (K) operators, which represent the average electron-electron repulsion and the antisymmetry of the wavefunction, respectively. It is this equation that is solved iteratively in the self-consistent field (SCF) procedure. The Fock matrix depends on the molecular

orbital coefficients (the solution itself), creating a circular dependency: the coefficients are needed to construct the Fock matrix, but solving the Fock matrix produces the coefficients. This self-consistent iteration continues until convergence, meaning the coefficients no longer change between iterations, hence the name "self-consistent field." (SCF). The variational principle ensures that the HF energy is an upper bound to the exact ground state energy, $E_{HF} = \langle \Phi_{HF} | \hat{H} | \Phi_{HF} \rangle \geq E_{exact}$ [12]. Which allows us to define that the correlation energy is negative: $E_{corr} = E_{exact} - E_{HF} < 0$. The resulting orbital energies ϵ_p provide insight into the electronic structure, with occupied orbitals typically having negative energies and virtual orbitals having positive energies, as per Koopmans' theorem [3].

Restricted and Unrestricted Formalism

Hartree-Fock theory varies in its treatment of spin, leading to two different formalisms: Restricted Hartree-Fock (RHF) and Unrestricted Hartree-Fock (UHF). Spin orbitals are constructed from spatial orbitals and spin functions:

$$\chi_p(\mathbf{r}, \sigma) = \phi_p(\mathbf{r})\sigma \quad (2.4)$$

where $\sigma \in \{\alpha, \beta\}$ denotes the spin function for spin-up and spin-down electrons, respectively.

In RHF, the same spatial orbital ϕ_p is used for both α and β electrons, expressed as the constraint $\phi_\alpha = \phi_\beta$. Consequently, each spatial orbital is doubly occupied by electrons of opposite spin, forming a single set of molecular orbitals. This constraint makes RHF suitable for closed-shell systems and ensures that the wavefunction is an eigenfunction of the total spin operator $\hat{S}^2$, keeping the total spin of the system fixed [3]. However, RHF fails to describe open-shell systems and bond dissociation correctly, as the spin-symmetry constraint cannot be satisfied in these cases [13].

In contrast, UHF allows different spatial orbitals for α and β spins ($\phi_\alpha \neq \phi_\beta$), making it applicable to open-shell systems [3]. This flexibility introduces spin contamination, where $\langle \hat{S}^2 \rangle$ deviates from the ideal value for a pure spin state [13]. In bond breaking and open-shell scenarios, UHF provides qualitatively correct descriptions where RHF fails [3]. For this reason, UHF is used throughout this thesis for generality, as it handles both closed-shell and open-shell systems effectively. While UHF incurs higher computational cost due to separate α and β orbital sets, incurring an additional degree of freedom, the scaling with the number of basis functions remains similar to RHF [3].

2.2.2 Møller-Plesset Perturbation Theory

Møller-Plesset perturbation theory (MPPT) treats electron correlation as a perturbation on top of the mean-field Fock operator. Truncated at second order (MP2), it is the simplest post-HF method that captures dynamical correlation at $\mathcal{O}(O^2V^2)$ computational cost, and furnishes the energy functional that OVOS seeks to optimise.

MP2 Correlation Energy

The MP2 correlation energy is derived by partitioning the electronic Hamiltonian as $\hat{H} = \hat{f} + \lambda\hat{V}$, where $\hat{f}$ is the Fock operator, which is diagonal in the canonical molecular

orbital basis, and $\hat{V}$ is the fluctuation potential that accounts for electron correlation [3]. The zeroth-order solution is the Hartree-Fock determinant $|\Phi_0\rangle$. Brillouin’s theorem states that singly-excited determinants have zero coupling to $|\Phi_0\rangle$, consequently, first-order perturbation theory contributes no singles terms [3, 13]. The first-order correction to the wavefunction is therefore a linear combination of doubly-excited determinants only. The associated MP1 amplitudes for these double excitations are:

$$t_{ij}^{ab} = -\frac{\langle ab||ij\rangle}{\varepsilon_a + \varepsilon_b - \varepsilon_i - \varepsilon_j} \quad (2.5)$$

where $\langle ab||ij\rangle$ is the antisymmetrised two-electron integral in physicists’ notation, and the denominator $\varepsilon_a + \varepsilon_b - \varepsilon_i - \varepsilon_j$ is the energy cost of exciting occupied orbitals i, j to virtual orbitals a, b . [3, 13].

The MP2 correlation energy over all occupied and virtual orbitals is:

$$E_{corr}^{(2)} = \frac{1}{4} \sum_{i,j,a,b} t_{ij}^{ab} \langle ij||ab\rangle \quad (2.6a)$$

$$= \sum_{i>j, a>b} t_{ij}^{ab} \langle ij||ab\rangle \quad (2.6b)$$

where i, j denotes occupied orbitals and virtual orbitals denoted by a, b [3]. In practice, the summation can be restricted to $i > j$ and $a > b$ to avoid double counting, leading to 2.6b. It is critical to make clear that both of these expressions sum over the entire orbital space, the full MP2 energy. In section 2.6, we will contrast this with the truncated MP2 energy used in OVOS, a constricted orbital space, reducing computational cost while retaining most correlation energy.

High-Order Amplitudes and T1 definition

In MP2, the MP1 amplitudes correspond to double electron excitations (T2 amplitudes). The T1 amplitudes, which represent single excitations, do not appear in the MP2 energy expression due to Brillouin’s theorem, which states that the matrix elements of the Hamiltonian between the Hartree-Fock determinant and singly-excited determinants vanish [3]. The presence of non-zero T1 amplitudes is an indicator that the reference determinant is not optimal, as it suggests that the orbitals can be rotated to lower the energy [11, 18]. In the context of coupled cluster theory, the stationarity condition $\nabla_R E(R) = 0$ defines optimal orbitals, which are formalized in quantum chemistry as Brueckner orbitals (named after Brueckner’s self-consistency principle) where the T1 amplitudes vanish [18]. The connection to Brueckner’s foundational work on self-consistency is discussed in [19]. This connection between orbital optimization and the nature of the excitation amplitudes is a key motivation for developing methods like OVOS that seek to optimize the virtual orbital space to minimize correlation energy [6].

MP2 Limitations and Applicability

MP2 is size-consistent[3], meaning that the correlation energy of non-interacting subsystems is additive 2.6b, which is a crucial property for accurately describing large systems and reaction energies [13]. However, MP2 is a single-reference method and

can fail in cases where the reference determinant is not a good approximation to the true wavefunction, such as in bond dissociation or transition states [3, 5]. In such cases, higher-order methods or multi-reference approaches may be necessary to capture the correct physics [1].

Computational Scaling Motivation

The $\mathcal{O}(O^2V^2)$ computational scaling of MP2 presents a significant barrier for large systems and higher-order methods [5, 13]. The Optimized Virtual Orbital Space (OVOS) method, detailed in Chapter 2, applies a strategic solution: by restricting the correlation energy calculation to a small, actively-correlated subset of virtual orbitals, the effective scaling can be reduced to $\mathcal{O}(N_{occ}^2 N_{virt}'^2)$ or better, where $N_{virt}' \ll N_{virt}$ [6]. This active space reduction is particularly powerful when combined with higher-order methods like CCSD(T), which scale as $\mathcal{O}(O^3V^4)$ or worse in the full space. For which V reflects the size of the virtual space, and reducing V to N_{virt}' can make previously intractable calculations feasible while retaining much of the correlation energy. This is a key motivation for the development of OVOS and its application to larger systems and more accurate methods.

2.3 Background for understanding OVOS

...

2.3.1 Orbital Rotations and Stationarity

Orbital optimization in post-HF methods relies on rotating the molecular orbitals to minimize a target energy functional, such as the MP2 correlation energy in the case of OVOS. This section introduces the mathematical framework for orbital rotations and the stationarity principle that defines optimal orbitals.

Partitioning and Index Notation

Orbitals in electronic structure theory are often partitioned into different subspaces to facilitate the formulation of post-HF methods. The Optimised Virtual Orbital Space (OVOS) method partitions the full spin-orbital basis into three non-overlapping subsets: occupied, active unoccupied (active virtual), and inactive unoccupied (inactive virtual), See Figure 2.3. The occupied orbitals, denoted by indices i, j, k, l , correspond to the spin-orbitals that are occupied in the Hartree-Fock reference state, and they are indexed from 0 to $n_{elec} - 1$, where n_{elec} is the number of electrons. The active virtual orbitals, denoted by indices a, b, c, d , are the subset of unoccupied orbitals that are optimised in the OVOS method, and they are indexed from n_{elec} to $n_{elec} + N_{virt}' - 1$, where N_{virt}' is the number of active virtuals and is a key parameter in the OVOS method. The inactive virtual orbitals, denoted by indices e, f, g, h , are the remaining unoccupied orbitals that are not explicitly rotated in the OVOS method; they define the complement space and are indexed from $n_{elec} + N_{virt}'$ to $n_{spin} - 1$, where n_{spin} is the total number of spin-orbitals. The OVOS rotation operates only in the block that couples the active virtuals with the inactive virtuals, while the occupied orbitals

remain fixed. The full set of virtual orbitals can be denoted as $a..e$, which includes all orbitals from n_{elec} to $n_{spin} - 1$ for a full MP2 calculation. This partitioning is illustrated in Figure 2.3, and the index notation is summarised in Table 2.1.




Figure 2.3: Partitioning of spin-orbital space for OVOS into occupied (blue), active virtual (purple), and inactive virtual subspaces (red). (1) No partitioning. (2) Occupied and unoccupied orbitals. (3) Active and inactive subsets.

Label	Indices	Orbital Space	Range	Count
Occupied	i, j, k, l	occ.	$[0, n_{elec})$	n_{elec}
Active Virtual	a, b, c, d	act. virt.	$[n_{elec}, n_{elec} + N'_{virt})$	N'_{virt}
Inactive Virtual	e, f, g, h	inact. virt.	$[n_{elec} + N'_{virt}, n_{spin})$	$N_{virt} - N'_{virt}$
All Virtual	$a..h$	full virt.	$[n_{elec}, n_{spin})$	N_{virt}

Table 2.1: Orbital subspace notation used throughout this thesis.

By this partitioning scheme, the OVOS method focuses on optimising the active virtual subspace to minimize the correlation energy, while the inactive virtuals provide a fixed reference that defines the complement space [6]. This truncation reduces the computational cost from $\mathcal{O}(O^2V^2)$ to $\mathcal{O}(O^2N'^2)$, where N' is the number of active virtuals, which is typically much smaller than the total number of virtuals $N' < V$. The choice of N' is a critical aspect of the OVOS method, as it involves a trade-off between accuracy and computational cost.

Orbital Rotation

Orbital rotations can be represented mathematically as unitary transformations of the molecular orbital coefficients [3]. A unitary transformation can be expressed as $U = e^R$, where R is an antisymmetric matrix that parameterizes the rotation. The new orbital coefficients after rotation are given by $C_{new} = C_{old}U$. The effect of this rotation on the one-electron integrals is given by $h_{new} = U^\dagger h_{old}U$, and for the two-electron integrals, the transformation is more complex due to their four-index nature, but it can be expressed as $g_{new} = U^\dagger g_{old}U$ in a suitable tensor contraction form. The Fock matrix also transforms under the same unitary rotation: $F_{new} = U^\dagger F_{old}U$. The use of unitary transformations ensures that the orthonormality of the orbitals is preserved during the optimization process [3].

The Stationarity Principle

The energy of the system can be viewed as a functional of the rotation parameters, $E(R)$. The gradient of the energy with respect to these parameters, $\nabla_R E(R)$, is known as the orbital gradient [3, 11]. The stationarity condition for optimal orbitals is defined by setting this gradient to zero: $\nabla_R E(R) = 0$. This condition implies that the energy is at a local extremum with respect to orbital rotations. The Hessian matrix, $\nabla_R^2 E(R)$, provides information about the curvature of the energy landscape around the stationary point and can be used to determine whether the stationary point is a minimum, maximum, or saddle point [13]. The connection between this stationarity principle and specific types of optimal orbitals, such as OO-MP2 and Brueckner orbitals, arises from the particular form of the energy functional being optimized [11, 18].

Optimization Strategies

Several optimization strategies can be employed to find the optimal orbitals that satisfy the stationarity condition. Newton's method is a common approach that uses both the gradient and the Hessian to achieve fast local convergence [13]. However, it can be sensitive to the choice of initial guess and may not always converge. Trust region methods address this issue by restricting the step size to a region where the model is expected to be accurate, improving robustness [13]. The Levenberg-Marquardt algorithm is a popular choice that interpolates between Newton's method and gradient descent, providing a balance between convergence speed and stability [13]. Convergence criteria typically involve thresholds for the norm of the gradient and changes in energy between iterations.

At convergence, when $\nabla_R E(R) = 0$, the stationarity condition becomes equivalent to the Brueckner condition $t_i^a = 0$ for all occupied-virtual pairs (i,a) [11, 18]. This fundamental connection shows that Brueckner orbitals are simply the stationary point of the orbital optimization for a chosen energy functional. The Brueckner orbitals are discussed in the following section.

2.3.2 The Second-Order Hylleraas Functional

OVOS finds its roots in minimising the second-order Hylleraas functional J_2 with respect to the choice of active virtual orbitals [6]. This section seeks to derive J_2 from the second-order Rayleigh-Schrödinger perturbation expression and to establish its connection to the MP2 correlation energy.

From Perturbation Theory to the Hylleraas Principle

We arrive at the Hylleraas functional by starting from the Rayleigh-Schrödinger (RS) perturbation theory expression for the second-order energy correction [3], starting by an expansion of the total energy:

$$E = E^{(0)} + E^{(1)} + E^{(2)} + \dots \quad (2.7)$$

where the zeroth-order energy $E^{(0)}$ is the energy of the zero-order Hamiltonian H_0 (The Fock Operator), and the first-order energy $E^{(1)}$ is the expectation value of

the correlation perturbation operator $V = H - H_0$ with respect to the zero-order wavefunction $\Psi^{(0)}$,

$$E^{(1)} = \langle \Psi^{(0)} | V | \Psi^{(0)} \rangle \quad (2.8)$$

The sum of zero-order and first-order energies, $E^{(0)} + E^{(1)}$, is the Hartree-Fock energy.

The second-order energy correction $E^{(2)}$ is the MP2 correlation energy, expressed in terms of the correlation perturbation operator $V = H - H_0$ and the first-order wavefunction $\Psi^{(1)}$ as:

$$E^{(2)} = \langle \Psi^{(0)} | V | \Psi^{(1)} \rangle \quad (2.9)$$

where $\Psi^{(0)}$ is the Hartree-Fock reference wavefunction. Using intermediate normalization, $\langle \Psi^{(0)} | \Psi^{(1)} \rangle = 0$, and the first-order Schrödinger equation:

$$H_0 | \Psi^{(1)} \rangle + V | \Psi^{(0)} \rangle = E^{(0)} | \Psi^{(1)} \rangle + E^{(1)} | \Psi^{(0)} \rangle \quad (2.10)$$

we can rearrange to express:

$$(H_0 - E^{(0)}) | \Psi^{(1)} \rangle = -(V - E^{(1)}) | \Psi^{(0)} \rangle \quad (2.11)$$

take the inner product with $\langle \Psi^{(1)} |$:

$$\langle \Psi^{(1)} | H_0 - E^{(0)} | \Psi^{(1)} \rangle = -\langle \Psi^{(1)} | V - E^{(1)} | \Psi^{(0)} \rangle \quad (2.12)$$

using, $\langle \Psi^{(1)} | E^{(1)} | \Psi^{(0)} \rangle = E^{(1)} \langle \Psi^{(1)} | \Psi^{(0)} \rangle = 0$, we can rewrite the second-order energy correction as:

$$E^{(2)} \leq \langle \Psi^{(1)} | H_0 - E^{(0)} | \Psi^{(1)} \rangle + 2 \langle \Psi^{(1)} | V - E^{(1)} | \Psi^{(0)} \rangle \quad (2.13)$$

The second-order Hylleraas functional J_2 is defined as the upper bound to the second-order energy correction, such that $E^{(2)} \leq J_2$ for any choice of $\Psi^{(1)}$. In other words, minimising J_2 with respect to $\Psi^{(1)}$ recovers the MP2 correlation energy, $E^{(2)} = \min_{\Psi^{(1)}} J_2$. Minimising by varying $\Psi^{(1)}$ leads to the stationary condition that $\Psi^{(1)}$ must satisfy the first-order Schrödinger equation,

$$\frac{\partial J_2}{\partial \Psi^{(1)}} = 2(H_0 - E^{(0)}) | \Psi^{(1)} \rangle + 2(V - E^{(1)}) | \Psi^{(0)} \rangle = 0 \quad (2.14)$$

making the ansatz that $\Psi^{(1)}$ is a linear combination of doubly-excited determinants, $\Psi^{(1)} = \sum_{i>j, a>b} t_{ij}^{ab} | \Phi_{ij}^{ab} \rangle$, and $\Psi^{(0)}$ is the HF reference determinant, we can write out the MP2 amplitudes t_{ij}^{ab} as:

$$t_{ij}^{ab} = \frac{\langle ij || ab \rangle}{\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b} \quad (2.15)$$

where $\langle ij || ab \rangle$ are the antisymmetrised two-electron integrals in physicists' notation, and ϵ_p are the orbital energies.

Pair Decomposition and Explicit Form of J_2

To apply the Hylleraas functional to orbital optimisation, we decompose J_2 into unique pairs of occupied orbital indices $i > j$, each pair involve orthogonal configurations, $\langle \Phi_{ij}^{ab} | \Phi_{kl}^{cd} \rangle = 0$ for $(i, j) \neq (k, l)$, which translate to the functional being additively separable, $J_2 = \sum_{i>j} J_{ij}^{(2)}$, each pair contributes:

$$J_{ij}^{(2)} = \langle \Psi_{ij}^{(1)} | H_0 - E^{(0)} | \Psi_{ij}^{(1)} \rangle + 2 \langle \Psi_{ij}^{(1)} | V - E^{(1)} | \Psi^{(0)} \rangle \quad (2.16)$$

where we moved the sum over $i > j$ inside the definition of $\Psi_{ij}^{(1)}$ out to allow writing the decomposed form. Furthermore, we can decompose $J_{ij}^{(2)}$ into two contributions: the expectation value of the zero-order Hamiltonian (first term) and the coupling to the reference state via the perturbation operator (second term).

The Zero-Order Term

The "diagonal" term takes form as we evaluate the expectation value:

$$\langle \Psi_{ij}^{(1)} | H_0 - E^{(0)} | \Psi_{ij}^{(1)} \rangle = \sum_{a>b, c>d} t_{ij}^{ab} t_{ij}^{cd} \langle \Phi_{ij}^{ab} | (H_0 - E^{(0)}) | \Phi_{ij}^{cd} \rangle \quad (2.17)$$

To evaluate, we recall that $H_0 = \sum_i f_{ii}$, where f_{pq} is the Fock matrix element. The first-order energy is $E^{(0)} = \sum_i \epsilon_i$, which represents the sum of occupied orbital energies. The matrix element of H_0 between the doubly-excited determinants can be evaluated to:

$$\langle \Phi_{ij}^{ab} | H_0 | \Phi_{ij}^{cd} \rangle = f_{ac} \delta_{bd} - f_{ad} \delta_{bc} + f_{bd} \delta_{ac} - f_{bc} \delta_{ad} \quad (2.18)$$

The four term Fock structure arise from the antisymmetrised electron exchange and orbital index rearrangements, $\langle \Phi_{ij}^{ab} | f_{pq} | \Phi_{ij}^{cd} \rangle$ is non-zero only when the orbital indices match, which is selected by the δ structure.

The matrix element of $E^{(0)}$ is:

$$\langle \Phi_{ij}^{ab} | E^{(0)} | \Phi_{ij}^{cd} \rangle = (\epsilon_i + \epsilon_j) (\delta_{ac} \delta_{bd} - \delta_{ad} \delta_{bc}) \quad (2.19)$$

combined, the "diagonal" term becomes:

$$\langle \Psi_{ij}^{(1)} | H_0 - E^{(0)} | \Psi_{ij}^{(1)} \rangle = \sum_{a>b, c>d} t_{ij}^{ab} t_{ij}^{cd} [(f_{ac} \delta_{bd} - f_{ad} \delta_{bc} + f_{bd} \delta_{ac} - f_{bc} \delta_{ad}) \quad (2.20)$$

$$- (\epsilon_i + \epsilon_j) (\delta_{ac} \delta_{bd} - \delta_{ad} \delta_{bc})] \quad (2.21)$$

Evaluation of the Perturbative Coupling

The "off-diagonal" term, consisting of the perturbation operator V and the first-order energy $E^{(1)}$, assesses the coupling between the reference determinant and the doubly-excited determinants. Evaluating the matrix element gives for each:

$$\langle \Psi_{ij}^{(1)} | V | \Psi^{(0)} \rangle = \sum_{a>b} t_{ij}^{ab} \langle ij || ab \rangle \quad (2.22)$$

and,

$$\langle \Psi_{ij}^{(1)} | E^{(1)} | \Psi^{(0)} \rangle = E^{(1)} \langle \Psi_{ij}^{(1)} | \Psi^{(0)} \rangle = 0 \quad (2.23)$$

therefore the "off-diagonal" term reduces to:

$$2\langle\Psi_{ij}^{(1)}|V - E^{(1)}|\Psi^{(0)}\rangle = 2\sum_{a>b}t_{ij}^{ab}\langle ij||ab\rangle \quad (2.24)$$

The final expression for $J_{ij}^{(2)}$ is then:

$$J_{ij}^{(2)} = \sum_{a>b,c>d}t_{ij}^{ab}t_{ij}^{cd}(f_{ac}\delta_{bd}-f_{ad}\delta_{bc}+f_{bd}\delta_{ac}-f_{bc}\delta_{ad}-(\epsilon_i+\epsilon_j)(\delta_{ac}\delta_{bd}-\delta_{ad}\delta_{bc}))+2\sum_{a>b}t_{ij}^{ab}\langle ij||ab\rangle \quad (2.25)$$

2.3.3 Orbital Rotation Parameterisation

Optimising orbitals in OVOS is achieved by applying a unitary transformation, rotating the virtual orbitals to maximise the correlation energy captured by the active space [6]. This section introduces the mathematical formalism of orbital rotations, parameterised by an antisymmetric matrix $\mathbf{R}$, and how the unitary transformation $\mathbf{U} = e^{\mathbf{R}}$ is constructed to mix only the active and inactive virtual orbitals.

Active and Inactive Virtual Subspaces

The antisymmetric matrix $\mathbf{R}$ is defined such that only the block corresponding to active-inactive virtual orbital mixing is non-zero. Specifically, R_{ae} is non-zero, while all other blocks (e.g., occupied-occupied R_{ij} , active-active R_{ab} , inactive-inactive R_{ef}) are zero.

We can restrict the orbital rotations to only mix active and inactive virtual orbitals because active-active rotations do not change the value of J_2 , as J_2 depends only on the choice of optimised orbitals and not on their specific linear combinations of basis states. In other words, any rotation within the active space leaves J_2 invariant:

$$J_2[\{\phi_a\}] = J_2[\{\sum_b U_{ab}\phi_b\}] \quad (2.26)$$

Inactive-inactive rotations do not affect the active space at all. We neglect all mixing with and between occupied orbitals, as it is not within the scope of the OVOS method.

We can write the full matrix $\mathbf{R}$ as:

$$\mathbf{R} = \begin{bmatrix} R_{ij} & R_{ia} & R_{ie} \\ R_{aj} & R_{ab} & R_{ae} \\ R_{ej} & R_{eb} & R_{ef} \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & R_{ae} \\ 0 & -R_{ae}^T & 0 \end{bmatrix} \quad (2.27)$$

where we relate the active-inactive block to its transpose by the antisymmetry condition, $R_{ea} = -R_{ae}$. The dimension of the non-zero block is $N' \times (N_{\text{unocc.}} - N')$.

Unitary Transformation via Matrix Exponential

The unitary transformation $\mathbf{U}$ that rotates the orbitals is given by the matrix exponential of $\mathbf{R}$:

$$\mathbf{U} = e^{\mathbf{R}} = \sum_{n=0}^{\infty} \frac{\mathbf{R}^n}{n!} \quad (2.28)$$

Unitary by construction, $U^\dagger U = I$, is guaranteed by the antisymmetry of R , $R = -R^\dagger$. Expanding the virtual orbitals to second order in R gives:

$$\phi'_a = \sum_b U_{ab} \phi_b + \sum_e U_{ae} \phi_e \quad (2.29)$$

we can express an active orbital after rotation as:

$$\phi'_a = \phi_a + \sum_e R_{ea} \phi_e - \frac{1}{2} \sum_{b,e} R_{ea} R_{eb} \phi_b + \dots \quad (2.30)$$

where the original active orbital is mixed with the inactive orbitals via the first-order term, and the second-order term accounts for the fact that the new active orbital is no longer orthogonal to the original active orbitals. These terms are redundant in the face of applying the full matrix exponential instead, as it handles all terms, the expansion allows for a tractable expression for the gradient and Hessian of J_2 with respect to the rotation parameters, which are derived in the following sections.

2.3.4 Gradient Expression

The descent direction for OVOS optimisation is given by the gradient of the Hylleraas functional J_2 with respect to the orbital rotation parameters R_{ae} [6]. This section derives the expression for the gradient, which consists of two terms: a direct contribution from the active-inactive mixing and an indirect contribution via the MP2 virtual-virtual density matrix D_{ab} .

MP2 Virtual-Virtual Density Matrix

The second-order density matrix, captures the correlation between pairs of active virtual orbitals, how strongly active virtual orbitals a and b appear together, in the MP1 wavefunction, is defined as:

$$D_{ab} = \sum_{i>j} \sum_c t_{ij}^{ac} t_{ij}^{bc} \quad (2.31)$$

where t_{ij}^{ac} are the MP2 amplitudes. The matrix D_{ab} is symmetric, $D_{ab} = D_{ba}$,

Gradient Expression

The gradient of J_2 with respect to the orbital rotation parameters R_{ae} is given by:

$$G_{ae} = \left. \frac{\partial J_2}{\partial R_{ae}} \right|_{R=0} \quad (2.32)$$

here we evaluate the gradient at $R = 0$ to obtain the expression in terms of the unrotated orbitals. This condition means the Fock matrix and integrals are evaluated in the original orbital basis, such that the orbitals are canonical in the unoccupied space. This means in the Hylleraas functional, the Fock matrix elements f_{ac} are diagonal in the virtual space, $f_{ac} = \epsilon_a \delta_{ac}$, meaning we can express the "zero-order"

term of J_2 in terms of the orbital energies:

$$\langle \Psi_{ij}^{(1)} | H_0 - E^{(0)} | \Psi_{ij}^{(1)} \rangle = \sum_{a>b, c>d} t_{ij}^{ab} t_{ij}^{cd} [(\epsilon_a \delta_{ac} \delta_{bd} - \epsilon_a \delta_{ad} \delta_{bc} + \epsilon_b \delta_{bd} \delta_{ac} - \epsilon_b \delta_{bc} \delta_{ad}) \quad (2.33)$$

$$- (\epsilon_i + \epsilon_j)(\delta_{ac} \delta_{bd} - \delta_{ad} \delta_{bc})] \quad (2.34)$$

$$= \sum_{a>b, c>d} t_{ij}^{ab} t_{ij}^{cd} [(\epsilon_a + \epsilon_b - \epsilon_i - \epsilon_j)(\delta_{ac} \delta_{bd} - \delta_{ad} \delta_{bc})] \quad (2.35)$$

using the definition of the MP2 amplitudes, we can express the "zero-order" term as:

$$\langle \Psi_{ij}^{(1)} | H_0 - E^{(0)} | \Psi_{ij}^{(1)} \rangle = - \sum_{a>b, c>d} t_{ij}^{cd} \langle ij || ab \rangle (\delta_{ac} \delta_{bd} - \delta_{ad} \delta_{bc}) \quad (2.36)$$

$$= - \sum_{a>b} t_{ij}^{ab} \langle ij || ab \rangle \quad (2.37)$$

as only the first deltas term contributes, by the summation convention. Recognize the resulting expression as the "perturbative coupling" term, therefore the J_2 functional can be expressed as:

$$J_2 = \sum_{i>j} \sum_{a>b} t_{ij}^{ab} \langle ij || ab \rangle = \sum_{i>j} \sum_{a>b} \langle ij || ab \rangle \langle ij || ab \rangle / (\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b) \quad (2.38)$$

Using the chain rule to evaluate the gradient, in regards to the two-electron integrals and the Fock matrix elements, we get the intermediate expression:

$$\frac{\partial J_2}{\partial R_{ae}} = \sum_{i>j} \sum_{a>b} 2 \langle ij || ab \rangle / (\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b) \frac{\partial \langle ij || ab \rangle}{\partial R_{ae}} \quad (2.39)$$

$$- \sum_{i>j} \sum_{a>b} (\langle ij || ab \rangle / (\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b))^2 \frac{\partial}{\partial R_{ae}} (\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b) \quad (2.40)$$

here each term is evaluated to:

$$\frac{\partial \langle ij || ab \rangle}{\partial R_{ae}} = \langle ij || eb \rangle \quad \text{and} \quad \frac{\partial}{\partial R_{ae}} (\epsilon_i + \epsilon_j - \epsilon_a - \epsilon_b) = -f_{ea} \quad (2.41)$$

arriving at the final expression for the gradient:

$$G_{ae} = \frac{\partial J_2}{\partial R_{ae}} = 2 \sum_{i>j} \sum_b t_{ij}^{ab} \langle ij || eb \rangle + 2 \sum_b D_{ab} f_{eb} \quad (2.42)$$

for which we have substituted the definition of the density matrix D_{ab} and renamed the indices for clarity. The first term captures the direct contribution of the active-inactive mixing to J_2 , while the second term captures the contribution via the density matrix, which accounts for how optimised orbitals are correlated with each other in the MP1 wavefunction. At convergence, the gradient must vanish, $G_{ae} = 0$ for all a, e , a stationarity condition that defines the optimal orbitals. The gradient is a matrix of dimension $(N_{\text{virt}} - N') \times N'$.

2.3.5 Hessian Expression

Taken the second derivative of the Hylleraas functional with respect to the orbital rotation parameters R_{ae} and R_{fb} , we arrive at the Hessian matrix:

$$H_{ae,bf} = \frac{\partial^2 J_2}{\partial R_{ae} \partial R_{bf}} \Big|_{R=0} \quad (2.43)$$

If the gradient captures the slope of the energy landscape, the Hessian captures the curvature, providing information about how the energy changes as we move in different directions in orbital space.

Adamowicz and Bartlett derive the full expression for the Hessian, which consists of four terms:

$$H_{ae,bf} = 2 \sum_{i>j} t_{ij}^{ab} \langle ij || ef \rangle - \sum_{i>j} \sum_c [t_{ij}^{ac} \langle ij || bc \rangle + t_{ij}^{bc} \langle ij || ca \rangle] \delta_{ef} \quad (2.44)$$

$$+ D_{ab}(f_{aa} - f_{bb}) \delta_{ef} + D_{ab} f_{ef} (1 - \delta_{ef}) \quad (2.45)$$

they note the active orbitals are canonical, and that the Hessian also shows pair separability [6].

Derivation of the Hessian Expression

We start from the MP2 gradient written with consistent dummy indices:

$$G_{ae} = 2 \sum_{i>j} \sum_c t_{ij}^{ac} \langle ij || ec \rangle + 2 \sum_c D_{ac} f_{ec} \quad (2.46)$$

where

$$D_{ac} = \sum_{i>j} \sum_d t_{ij}^{ad} t_{ij}^{cd}, \quad f_{ec} = \epsilon_e \delta_{ec} \quad (2.47)$$

The Hessian is defined as:

$$H_{ea,fb} = \frac{\partial G_{ae}}{\partial R_{bf}} \quad (2.48)$$

We use the following standard orbital-rotation derivatives.

Two-electron integrals:

$$\frac{\partial \langle ij || ec \rangle}{\partial R_{bf}} = \delta_{cb} \langle ij || ef \rangle + \delta_{eb} \langle ij || fc \rangle - \delta_{cf} \langle ij || eb \rangle - \delta_{ef} \langle ij || bc \rangle \quad (2.49)$$

Fock matrix (canonical orbitals):

$$\frac{\partial f_{ec}}{\partial R_{bf}} = \delta_{cb} f_{ef} + \delta_{eb} f_{fc} - \delta_{cf} f_{eb} - \delta_{ef} f_{bc} \quad (2.50)$$

Amplitudes: Rather than expanding the full quotient rule, we use:

$$t_{ij}^{ac} = \frac{\langle ij || ac \rangle}{\Delta_{ijac}} \quad (2.51)$$

and differentiate formally, allowing cancellation later:

$$\frac{\partial t_{ij}^{ac}}{\partial R_{bf}} = \frac{1}{\Delta_{ijac}} \frac{\partial \langle ij || ac \rangle}{\partial R_{bf}} - t_{ij}^{ac} \frac{1}{\Delta_{ijac}} \frac{\partial \Delta_{ijac}}{\partial R_{bf}} \quad (2.52)$$

Step 1: Differentiate First Gradient Term

Define:

$$A = \frac{\partial}{\partial R_{bf}} \left(2 \sum_{i>j} \sum_c t_{ij}^{ac} \langle ij || ec \rangle \right) \quad (2.53)$$

Apply product rule:

$$A = 2 \sum_{i>j} \sum_c \left(\frac{\partial t_{ij}^{ac}}{\partial R_{bf}} \langle ij || ec \rangle + t_{ij}^{ac} \frac{\partial \langle ij || ec \rangle}{\partial R_{bf}} \right) \quad (2.54)$$

Integral derivative contribution

Insert the integral derivative:

$$\sum_c t_{ij}^{ac} \frac{\partial \langle ij || ec \rangle}{\partial R_{bf}} = \sum_c t_{ij}^{ac} (\delta_{cb} \langle ij || ef \rangle + \delta_{eb} \langle ij || fc \rangle - \delta_{cf} \langle ij || eb \rangle - \delta_{ef} \langle ij || bc \rangle) \quad (2.55)$$

Collapse sums:

$$= t_{ij}^{ab} \langle ij || ef \rangle + \delta_{eb} \sum_c t_{ij}^{ac} \langle ij || fc \rangle - t_{ij}^{af} \langle ij || eb \rangle - \delta_{ef} \sum_c t_{ij}^{ac} \langle ij || bc \rangle \quad (2.56)$$

Thus:

$$A_{\text{int}} = 2 \sum_{i>j} \left[t_{ij}^{ab} \langle ij || ef \rangle - t_{ij}^{af} \langle ij || eb \rangle \right] + 2\delta_{eb} \sum_{i>j} \sum_c t_{ij}^{ac} \langle ij || fc \rangle - 2\delta_{ef} \sum_{i>j} \sum_c t_{ij}^{ac} \langle ij || bc \rangle \quad (2.57)$$

Amplitude derivative contribution

The amplitude derivative produces terms of the same structure but with permuted indices. After relabeling dummy indices and combining symmetric contributions, these terms cancel the δ_{eb} terms and generate the symmetric partner of the δ_{ef} term:

$$-2\delta_{ef} \sum_{i>j} \sum_c t_{ij}^{bc} \langle ij || ca \rangle \quad (2.58)$$

Final result for A

$$A = 2 \sum_{i>j} t_{ij}^{ab} \langle ij || ef \rangle - \delta_{ef} \sum_{i>j} \sum_c \left[t_{ij}^{ac} \langle ij || bc \rangle + t_{ij}^{bc} \langle ij || ca \rangle \right] \quad (2.59)$$

Step 2: Differentiate Second Gradient Term

Define:

$$B = \frac{\partial}{\partial R_{bf}} \left(2 \sum_c D_{ac} f_{ec} \right) \quad (2.60)$$

$$B = 2 \sum_c \left(\frac{\partial D_{ac}}{\partial R_{bf}} f_{ec} + D_{ac} \frac{\partial f_{ec}}{\partial R_{bf}} \right) \quad (2.61)$$

Fock derivative term

$$\sum_c D_{ac} \frac{\partial f_{ec}}{\partial R_{bf}} = \sum_c D_{ac} (\delta_{cb} f_{ef} - \delta_{ef} f_{bc}) \quad (2.62)$$

$$= D_{ab} f_{ef} - \delta_{ef} \sum_c D_{ac} f_{bc} \quad (2.63)$$

Density derivative term

After inserting amplitude derivatives and simplifying (using symmetry and relabeling), this yields:

$$\delta_{ef} D_{ab} (f_{aa} - f_{bb}) \quad (2.64)$$

Final result for B

$$B = D_{ab} (f_{aa} - f_{bb}) \delta_{ef} + D_{ab} f_{ef} (1 - \delta_{ef}) \quad (2.65)$$

Final Hessian

$$H_{ea,fb} = 2 \sum_{i>j} t_{ij}^{ab} \langle ij || ef \rangle \quad (2.66)$$

$$- \delta_{ef} \sum_{i>j} \sum_c \left[t_{ij}^{ac} \langle ij || bc \rangle + t_{ij}^{bc} \langle ij || ca \rangle \right] \quad (2.67)$$

$$+ D_{ab} (f_{aa} - f_{bb}) \delta_{ef} + D_{ab} f_{ef} (1 - \delta_{ef}) \quad (2.68)$$

Block Structure of the Hessian

We see a clear simplification of the Hessian if $\delta_{ef} = 0$:

$$H_{ae,bf} = 2 \sum_{i>j} t_{ij}^{ab} \langle ij || ef \rangle + D_{ab} f_{ef} \quad (2.69)$$

such that two orbitals a, b would not rotate to the same orbitals, on the other hand if $\delta_{ef} = 1$:

$$H_{ae,be} = 2 \sum_{i>j} t_{ij}^{ab} \langle ij || ee \rangle - \sum_{i>j} \sum_c \left[t_{ij}^{ac} \langle ij || bc \rangle + t_{ij}^{bc} \langle ij || ca \rangle \right] + D_{ab} (f_{aa} - f_{bb}) \quad (2.70)$$

such that two orbitals a, b would rotate to the same orbital. This translate to a block structure within the Hessian, where each block corresponds to a specific inactive orbital e . Each block is a square matrix of size $N' \times N'$, where N' is the number of optimised orbitals. The off-diagonal blocks (where $e \neq f$) contain terms that couple different inactive orbitals, but these are typically smaller in magnitude than the diagonal blocks. This block structure approximation allows for efficient inversion of the Hessian, as each block can be inverted separately, reducing the computational cost from inverting a large matrix to inverting several smaller matrices.

Negative Eigenvalues and Level-Shift Stabilisation

As the Hessian expresses the curvature of the energy landscape, it may become indefinite (i.e., have negative eigenvalues) near saddle points or in cases with small virtual orbital energy gaps. This can lead to divergence or oscillations in the Newton-Raphson optimization. To address this, a level-shift strategy can be employed: if any diagonal block of the Hessian has negative eigenvalues, a positive shift λ is added to the diagonal elements, effectively replacing $\mathbf{H}$ with $\mathbf{H} + \lambda\mathbf{I}$. The choice of λ can be empirical or based on a dynamic threshold related to the smallest eigenvalue. This ensures that the modified Hessian is positive-definite, allowing for stable convergence towards a local minimum. Adamowicz & Bartlett does not discuss this strategy. It is observed that the block approximation does get rid of most negative eigenvalues, but some may still remain, like which may be in cases with small virtual gaps $\epsilon_a - \epsilon_i$ from basis sets with diffuse functions.

2.3.6 Newton-Raphson

Building off the gradient and Hessian of J_2 an orbital rotation update $\Delta\mathbf{R}$ can be computed using the Newton-Raphson (NR) method, which iteratively refines the orbital coefficients to minimize J_2 . The NR update is given by:

$$\mathbf{R} = -\mathbf{G} \cdot \mathbf{H}^{-1} \quad (2.71)$$

which is equivalent to:

$$\mathbf{H}\mathbf{R} = -\mathbf{G} \quad (2.72)$$

where $\mathbf{G}$ is the gradient vector and $\mathbf{H}$ is the Hessian matrix. The NR method has a quadratic convergence rate near the optimum, meaning that the error decreases quadratically with each iteration. In other words, if the current solution is close to the optimal orbitals, the next iteration will be significantly closer, leading to rapid convergence. However, if the energy landscape is flat or has multiple minima, convergence can be slow or may require damping of the update step. The block-diagonal structure of the Hessian allows for efficient inversion, as each block can be treated independently, which is equivalent to solving a reduced linear equation in each inactive-orbital block. To ensure numerical stability, especially in cases where the Hessian may have negative eigenvalues, a level-shift strategy can be applied, as discussed in the previous section. The convergence criterion for the NR iterations is typically defined as $|J_2(n) - J_2(n-1)| < 10^{-8}$, indicating that the change in the Hylleraas functional between iterations is sufficiently small to consider the solution converged.

2.3.7 Orbital Canonicalisation

After a unitary rotation, the virtual Fock block F_{ab}^{active} is generally non-diagonal, meaning the active virtual orbitals are no longer canonical [3]. To restore the canonical form, we compute the active Fock block in the rotated basis:

$$F_{ab} = \mathbf{C}'^\dagger F^{\text{AO}} \mathbf{C}' \quad (2.73)$$

where $\mathbf{C}'$ are the coefficients of the virtual orbitals after rotation. Diagonalising F_{ab} yields new canonical virtual orbitals with well-defined orbital energies ϵ_a from the

eigenvalues. This step is crucial for the next iteration of MP2 amplitude calculations, as it ensures that the denominators in the amplitude expressions are based on canonical orbital energies, simplifying the computations and improving numerical stability. It is important to note that only the virtual orbitals are rotated and re-canonicalised at each OVOS iteration; the occupied orbitals remain unchanged throughout the procedure.

With canonical rotated orbitals, the MP2 energy for the new active space can be evaluated, and the energy recovery fraction can be defined to quantify the performance of the OVOS procedure.

2.3.8 Computational Cost and Scaling

A central motivation for OVOS is the potential for significant computational speedup by contracting the orbital space to a smaller active orbital set [6]. To quantify this, we analyse the computational cost of each step in the OVOS iteration and compare it to the cost of traditional full-space correlation methods.

Per-Iteration Cost of OVOS Optimisation

Defining the cost of an action, requires introducing the Big O notation, which describes the asymptotic scaling of an algorithm with respect to the size of the input. In the context of OVOS, the relevant input sizes are the number of occupied orbitals (N_{occ}), the number of inactive virtual orbitals (N_{inact}), the number of optimized orbitals (N'), and the number of atomic orbitals (n).

The computational cost of one OVOS iteration can be broken down by the steps involved, see Section 4.1 for an algorithm summary, where the cost are:

- Integral transformation: $O(n^4 + N_{occ}^2 N_{virt}^2)$ for transforming two-electron integrals to the MO basis.
- MP2 amplitude evaluation: $O(N_{occ}^2 N'^2)$ for computing the amplitudes in the active space.
- MP2 energy: $O(N_{occ}^2 N'^4)$ for evaluating the MP2 energy with the active space amplitudes.
- Density matrix D_{ab} : $O(N_{occ}^2 N'^3)$ for computing the virtual-virtual density matrix.
- Gradient G_{ea} : $O(N_{occ}^2 N' N_{inact})$ for computing the gradient with respect to orbital rotations.
- Hessian $H_{ea,fb}$: $O(N_{occ}^2 N_{inact}^2 N'^2)$ for computing the Hessian in block-diagonal form.
- Block diagonal: $O(N_{inact} N'^2)$ for extracting the block-diagonal part of the Hessian.
- Newton-Raphson step: $O(N_{inact} N'^3)$ inverting the block-diagonal Hessian.
- Orbital rotation: $O((N_{occ} + N_{inact} + N')^3)$ for computing the unitary rotation matrix via matrix exponential.

- Canonicalisation: $O(N'^3)$ for diagonalising the active Fock block to obtain canonical orbitals.

Overall, the per-iteration cost is dominated differently for a few cases, summarized in the following table: This reflects the trade-off in OVOS: smaller active spaces reduce

Active Space	$r = N'/N_{\text{virt}}$	Dominant Step	Cost Scaling	Relative Cost
Small	~ 0.1	Orbital Rotation	$O(N_{\text{virt}}^3)$	Fastest, but minimal speedup
Medium	$\sim 0.3\text{--}0.5$	Newton Solver*	$O((N_{\text{inact}}N')^3)$	Optimal window
Large	$\sim 0.7\text{--}0.9$	MP2 Energy	$O(N_{\text{occ}}^2N'^4)$	Diminishing returns

Table 2.2: Dominant computational steps in OVOS iterations as a function of the active space size N' . *MP2 close 2nd.

the cost of MP2 energy and gradient evaluations but increase the relative cost of orbital rotations and Hessian inversions (number of blocks). Larger active spaces reduce the cost of orbital rotations but increase the cost of MP2 energy evaluations. The optimal window for OVOS is typically in the medium range, where the cost of the Newton step is manageable and the speedup from reduced MP2 evaluations is significant.

Taken the size of the basis set into account, translating to the number of virtual orbitals, the per-iteration cost of OVOS is dominated by the MP2 energy evaluation step, which scales as $O(N_{\text{occ}}^2N'^4)$. However, since $N' < N_{\text{virt}}$, this represents a significant reduction compared to the full-space MP2 cost as the basis set size increases, beating out the cost of orbital rotations and Hessian inversions.

Comparison: Full Virtual Space vs OVOS-Reduced

In a traditional full-space MP2 calculation, the cost scales as $O(N_{\text{occ}}^2N_{\text{virt}}^4)$ due to the need to consider all virtual orbitals in the amplitude calculations

Downstream Speedup Factors

In turn, how well would the reduced active space from OVOS speed up downstream methods like VQE. VQE scaling is dominated by the number of virtual orbitals, as the number of qubits and circuit depth scale with the size of the active space. The cost of VQE with a UtUPS ansatz scales as $O(N_{\text{qubit}})$, so the speedup from OVOS is the percentage of lower virtual space size, $O(N_{\text{virt}}/N')$. An advantage of OVOS might not be the recovered MP2 energy for a system, but the corresponding reduced virtual orbitals, as a trial state.

Breakeven Analysis and Practical Realisation

Despite a speedup seeming good, it would come at a cost of recovered correlation energy. The breakeven point for OVOS depends on the desired accuracy in the correlation energy and the size of the system. For small systems with few virtual orbitals, the overhead of OVOS iterations may not be justified, as the full-space MP2 calculation is already manageable. However, for larger systems with many virtual orbitals, OVOS can provide significant speedup while still recovering a substantial

fraction of the correlation energy, as will be discussed. The practical breakeven point will differ from performance benchmarks, as it depends on the specific system, basis set, and computational resources available.

2.4 Relevance to Quantum Computing

...

2.4.1 Brueckner Orbitals and Connection to OVOS

A special case of orbital optimisation in post-HF methods is the concept of Brueckner orbitals, which are defined as the orbitals that diagonalise the single-particle Green's function [19].

The single-particle Green's function is an advanced quantum mechanical object that encodes off-diagonal elements of the density matrix corrections from correlation [3]. Diagonalizing this function means the correlation-induced density matrix has minimal off-diagonal elements or, equivalently, the orbitals for which the T1 coupled-cluster amplitudes vanish for all occupied-virtual pairs, $t_{ia}^1 = 0$. Brueckner orbitals maximise the overlap of the Hartree-Fock determinant with the exact ground state wavefunction (FCI), $S = |\langle \Phi_{\text{HF}} | \Phi_{\text{exact}} \rangle|^2$, making Brueckner orbitals an expected optimal choice for reference orbitals in correlated calculations [18]. In practice, Brueckner orbitals are computed using the Brueckner-CCD (B-CCD) method [18]. This method iteratively solves the coupled-cluster doubles equations while simultaneously rotating the orbitals to eliminate the T1 amplitudes. The resulting Brueckner orbitals are often found to provide a better starting point for correlation treatments than the canonical HF orbitals, especially in cases where the HF reference is a poor approximation to the true ground state [18, 20]. The presence of non-zero T1 amplitudes in the HF reference indicates that the density matrix obtained from correlation differs from the HF density, leading to computational effort being wasted on single excitations [11].

Connection to Orbital-Optimised MP2 methods

Brueckner orbitals are closely related to the concept of orbital-optimised MP2 (OO-MP2) [11], which variationally minimises the MP2 energy with respect to all orbital rotations. At convergence, the gradient of the MP2 energy with respect to orbital rotations vanishes, which is the same stationarity condition that defines Brueckner orbitals at the MP2 level [11]. Therefore, OO-MP2 orbitals can be considered as the MP2-level Brueckner orbitals.

We can also draw a connection to the OVOS method, which minimises the MP2 energy in a restricted active virtual space [6]. Conceptually, this is the same orbital stationarity condition as for Brueckner orbitals, but applied only to the active block of virtual orbitals. The OVOS optimised molecular orbitals are hence expected to approximate Brueckner orbitals within the truncated active virtual space. This is a key novel claim of this thesis, and will be explored in detail in the Theory chapter Section 3.9, where we will investigate the extent to which OVOS orbitals approximate Brueckner orbitals and the implications for the performance of the VQE method when using OVOS orbitals as a reference state.

More precisely, the OVOS procedure enforces the stationary point condition $\nabla_R E(R) = 0$ for all active-inactive orbital rotations R_{ae} (where $a \in$ active virtuals, $e \in$ inactive virtuals). This stationarity condition is mathematically identical to enforcing the Brueckner condition $t_i^a = 0$ within the truncated active virtual subspace. In other words: OVOS-optimized orbitals ARE the Brueckner orbitals of the truncated orbital space. This insight has direct implications for higher-order correlation methods: just as Brueckner-CCD orbitals are known to dramatically reduce T1 and T2 amplitudes compared to HF, we expect OVOS orbitals to yield smaller T1 amplitudes in subsequent coupled-cluster calculations [18, 20].

OVOS Orbitals as Approximate Brueckner Orbitals

OVOS orbitals satisfy a stationarity condition with respect to orbital rotations between the active and inactive virtual spaces. Specifically, at convergence of the OVOS procedure, the gradient of the Hylleraas functional with respect to these rotations vanishes:

$$\frac{\partial J_2}{\partial R_{ae}} = G_{ae} = 0 \quad \forall a, e \quad (2.74)$$

This condition implies that the MP2 energy is stationary with respect to orbital rotations that mix the active virtual orbitals (those included in the OVOS optimization) and the inactive virtual orbitals (those excluded from the active space). In the context of orbital-optimized MP2 (OO-MP2), stationarity of the energy with respect to all orbital rotations (including occupied-virtual) leads to the condition that the single-excitation amplitudes T_1 vanish, which defines Brueckner orbitals.

We express the single-excitation amplitudes T_1 as the coupling between the reference determinant and singly excited determinants:

$$T_1 = \sum_{i,a} t_i^a |\Phi_i^a\rangle \quad (2.75)$$

where t_i^a are the single-excitation amplitudes, $|\Phi_i^a\rangle$ are the singly excited determinants obtained by replacing an occupied orbital i with a virtual orbital a .

While OVOS does not enforce stationarity with respect to occupied-virtual rotations, it achieves stationarity within the truncated active virtual space, effectively making the OVOS orbitals the Brueckner orbitals of that active space. This connection has important implications for the quality of the reference state and the expected reduction of T_1 amplitudes in subsequent correlation treatments.

2.4.2 Quantum Computing and Orbital Space

Quantum computing is an emerging paradigm that has the potential to revolutionise computational chemistry by enabling the simulation of quantum systems that are intractable for classical computers. Bridging from classical electronic structure methods to quantum algorithms requires careful consideration of how molecular orbitals are represented and utilised in the quantum computational framework.

Second Quantisation and Qubit Mapping

In the context of quantum chemistry, the electronic structure problem is often formulated in second quantisation, where the Hamiltonian is expressed in terms of creation

and annihilation operators acting on a reference state [3]. To implement quantum algorithms on a quantum computer, the fermionic operators must be mapped to qubit operators using transformations such as the Jordan-Wigner or Bravyi-Kitaev mappings [10]. The second quantised electronic Hamiltonian can be written as [5]:

$$\hat{H} = \sum_{pq} h_{pq} a_p^\dagger a_q + \frac{1}{2} \sum_{pqrs} g_{pqrs} a_p^\dagger a_q^\dagger a_r a_s, \quad a_p^\dagger a_q \rightarrow \mathcal{O}_{pq}, \quad a_p^\dagger a_q^\dagger a_r a_s \rightarrow \mathcal{O}_{pqrs}$$

This mapping allows us to represent the electronic Hamiltonian as a sum of Pauli operators, which can then be implemented as quantum circuits [5]. The choice of molecular orbitals used in the second quantised Hamiltonian can significantly impact the efficiency of the quantum algorithm, as it affects the structure of the Hamiltonian and the overlap of the reference state with the true ground state. In particular, using OVOS-optimised orbitals as the basis for the second quantised Hamiltonian might lead to a more compact representation and improved convergence properties in quantum algorithms such as the VQE [5, 6].

The Variational Quantum Eigensolver

The VQE [4] is a hybrid quantum-classical algorithm designed to find the ground-state energy of a quantum system [5]. VQE uses a parameterised quantum circuit, known as an ansatz, to prepare trial wavefunctions on a quantum computer. A classical outer loop optimises the parameters of the ansatz to minimise the energy expectation value, which is measured on the quantum device [4, 5].

INSERT DIAGRAM of vqe workflow here (see notes for details)

The choice of reference state in VQE is crucial for the efficiency and accuracy of the algorithm, and OVOS orbitals offer a promising avenue for improving the performance of VQE by providing a better initial guess for the ground state wavefunction [5]. Leaving the choice of ansatz, while Unitary Coupled Cluster Singles and Doubles (UCCSD) is a boardly used chemically-motivated ansatz in VQE [4, 5], the Unrestricted tiled Unitary Product State (UtUPS) ansatz for its simplified structure and choosen to focus on the impact of the reference state. The UtUPS ansatz is constructed as $|\Psi(\theta)\rangle = e^{T(\theta)-T^\dagger(\theta)}|\Phi_{\text{ref}}\rangle$, where $|\Phi_{\text{ref}}\rangle$ is a reference state, typically the Hartree-Fock determinant, and the Θ 's are parameters, tuning the ansatz [4]. The quality of the reference state significantly impacts the performance of VQE, as a reference with better overlap with the true ground state can lead to faster convergence and more compact orbital spaces [5].

OVOS Molecular Orbitals as a Trial State for VQE

We had VQE as a potential downstream application of OVOS, as the reduced active virtual space directly translates to fewer qubits and shallower circuits. VQE pivots us from classical quantum chemistry to quantum computing, where the quality of the trial state is crucial for convergence and accuracy. This section motivates why OVOS trial states are beneficial for variational quantum algorithms on NISQ devices.

Reference-State Overlap

To quantify the quality of the OVOS orbitals as a trial state for VQE, we can consider the overlap (or fidelity) between the reference state constructed from the OVOS

orbitals and the Full Configuration Interaction (FCI) ground state. The fidelity F is defined as:

$$F = |\langle \Phi_{ref} | \Psi_{FCI} \rangle|^2 \quad (2.76)$$

where Φ_{ref} is the reference state built from the OVOS orbitals and Ψ_{FCI} is the FCI wavefunction expressed in the natural orbital basis.

The claim is that $F_{OVOS} \geq F_{HF}$ because OVOS orbitals are optimized to minimize the MP2 energy, which should yield a reference state that has a higher overlap with the correlated wavefunction compared to the Hartree-Fock orbitals. This improved overlap can lead to faster convergence in VQE, as the initial state is closer to the true ground state, and can also reduce the required circuit depth to achieve a given accuracy.

Trial states

To evaluate how well OVOS orbitals will do, we formulate the following candidates: UMP2 natural orbitals, and canonical UHF orbitals. The former are the natural orbitals obtained from a UMP2 calculation, which are known to provide a good approximation to the correlated wavefunction and often yield improved convergence in VQE, formulated in PySCF [9]. They formulate we arrive at the natural orbitals by diagonalizing the one-particle density matrix constructed from the UMP2 amplitudes:

$$D_{pq} = \sum_{ijab} t_{ij}^{ab} \langle pq | ij \rangle \quad (2.77)$$

where t_{ij}^{ab} are the UMP2 amplitudes and $\langle pq | ij \rangle$ are the two-electron integrals. Diagonalizing D_{pq} yields the natural orbitals:

$$SD_{pq}SC = SCn \quad (2.78)$$

where S is the overlap matrix, C are the coefficients of the natural orbitals, and n are the natural orbital occupation numbers. The natural orbitals are expected to have a higher overlap with the FCI ground state compared to the canonical UHF orbitals, which are obtained from a standard Hartree-Fock calculation and may not capture correlation effects as effectively.

Reduction of T1 Amplitudes and Circuit Depth

Within VQE the reduction of T_1 amplitudes in the OVOS orbitals implies that the single-excitation contributions to the wavefunction are minimized. This has significant implications for the choice of ansatz in VQE, as it allows for the use of simple ansatz dropping single excitations (e.g., UCCD instead of UCCSD), which can significantly reduce the circuit depth and complexity. The reduction in T_1 amplitudes might reflect a good approximation to the reference state, which can lead to faster convergence and improved accuracy in VQE calculations.

Chapter 3

OVOS Implementation

Here in this chapter we build on the theoretical foundation laid out in Chapter 2 to describe the practical implementation of the OVOS algorithm. The goal is to provide a clear and detailed overview of the algorithmic steps, the computational strategies employed, and the choices made in the implementation. This chapter serves as a bridge between the abstract theoretical concepts and the concrete computational details, enabling readers to understand how the OVOS method is realised and how it can be reproduced or extended in future work.

3.1 Algorithm

Here we outline the full OVOS algorithm as implemented in this work, following the general structure proposed by Adamowicz & Bartlett [6]. The algorithm consists of an initialisation phase followed by an iterative loop that optimises the active virtual orbitals to maximise the MP2 correlation energy. The key steps are summarised as pseudocode in Algorithm 1, which serves as a roadmap for the detailed descriptions in the subsequent sections.

Algorithm 1 OVOS Iterative Optimisation

- 1: Perform SCF calculation to obtain initial MO coefficients $\mathbf{C}$
 - 2: Transform AO integrals to MO basis: $\langle ij|ab\rangle$
 - 3: Select initial active virtual space
 - 4: **repeat**
 - 5: Compute MP1 amplitudes t_{ij}^{ab} and MP2 energy J_2
 - 6: Assemble gradient $\mathbf{G}$ and Hessian $\mathbf{H}$
 - 7: Solve Newton-Raphson equation: $\Delta\mathbf{R} = -\mathbf{H}^{-1}\mathbf{G}$
 - 8: Generate unitary matrix $\mathbf{U} = \exp(\mathbf{R})$
 - 9: Update MO coefficients: $\mathbf{C} \leftarrow \mathbf{C}\mathbf{U}$
 - 10: Canonicalise active Fock block
 - 11: Evaluate new J_2
 - 12: **until** $|J_2^{(n)} - J_2^{(n-1)}| < \epsilon_{\text{conv}}$
-

Below we provide a detailed description of each step in the algorithm.

3.1.1 Integral Transformation

AO-to-MO Transformation

Spin-Orbital Construction

Notation Convention: Chemists' vs. Physicists' Integrals

3.1.2 Initial Active Space Selection

3.1.3 MP1 Amplitudes and MP2 Energy

3.1.4 Gradient and Hessian Assembly

Block Structure and Efficient Inversion

Level-Shift Stabilisation

3.1.5 Orbital Rotation and Update

3.1.6 Convergence and Iteration Control

3.2 Initialisation Strategies

3.2.1 Default SCF Initialisation

3.2.2 Warm-Start from Previous Orbital Set

3.2.3 Random Unitary Sampling

Chapter 4

Computational Detail

In this chapter, we provide a comprehensive account of the computational details underlying all calculations presented in this thesis. This includes the software and libraries used, the hardware environment, the molecular geometries and basis sets employed, the settings for the reference SCF calculations, and the specific hyperparameters of the OVOS optimisation. The goal is to ensure that all results are fully reproducible by any reader with access to similar computational resources.

4.0.1 Software and Libraries

In this section we detail the software environment used for all calculations in this thesis. The primary quantum chemistry calculations were performed using PySCF version 2.10 [9], and SlowQuant branch Unrestricted VQE [21]. Numerical operations were carried out using NumPy version 1.26.4 and SciPy version 1.17, a complete list of which is provided in Table 4.1.

Table 4.1: Software packages used in this work.

Package	Version	Purpose	Citation
PySCF	2.10.0	Quantum chemistry calculations	...
SlowQuant	A.B.C	Alternative quantum chemistry package	[21]
NumPy	1.26.4	Numerical operations	...
SciPy	1.17.0	Scientific computing	...

The OVOS method was implemented as a standalone Python class, with source code provided in Appendix C. All calculations were executed in a single-threaded mode to ensure reproducibility, with the environment variables `OPENBLAS_NUM_THREADS` and `OMP_NUM_THREADS` set to 1. The operating system used was Windows in Microsoft Visual Studio Code’s WSL ubuntu virtual machine, as PySCF needs a Unix-like environment, and the Python distribution was standard Python version 3.11.

The hardware environment on which all calculations were run is described in the following section.

4.0.2 Molecular Geometries

The thesis focuses on a test set of nine molecules that represent a variety of electronic structures, all abundantly used in quantum chemistry benchmarks, and found in nature. These include diatomic molecules (CO, HF), small polyatomic molecules (H₂O, NH₃), and a simple open-shell molecule (LiH). All geometries were kept fixed throughout the calculations, with no geometry optimisation performed. The specific coordinates used for each molecule are listed in Table 4.2, along with their spin multiplicity, charge, and

point group symmetry. The bond lengths and angles were taken from experimental values reported in the literature, ensuring that the test set is representative of real molecular systems. The coordinate unit used is Ångström.

Table 4.2: Molecular geometries and electronic structure parameters used in this work.

Molecule	Geometry (Å)
LiH	$r(\text{Li} - \text{H}) = 1.595$
H ₂ O	$r = 0.957, \angle = 104.5^\circ$
CO	$r(\text{C} - \text{O}) = 1.128$
HF	$r(\text{H} - \text{F}) = 0.917$
NH ₃	$r(\text{N} - \text{H}) = 1.012$

The basis-set-specific orbital space dimensions for each molecule are given in the following section, along with a description of the basis sets employed.

4.0.3 Basis Sets

The choice of basis set determines the size of the virtual orbital space and hence the difficulty and accuracy of the OVOS calculation. This section describes the basis sets employed and provides the resulting orbital space dimensions for each molecule. See Basis Set Exchange [22] for the source of all basis sets used in this work.

6-31G: A split-valence double-zeta basis set developed by Pople et. al. [14–16]. It consists of a minimal basis for core electrons and a double-zeta basis for valence electrons, allowing for improved flexibility in describing electron correlation compared to minimal basis sets.

cc-pVDZ: The correlation-consistent polarized valence double-zeta basis set developed by Dunning [17]. It includes polarization functions that allow for a better description of electron correlation effects, particularly important for accurate energy calculations.

4.1 OVOS Calculation Settings

For all calculations, the molecule is initialized using PySCF with `gto.Mole()` and the Hartree-Fock reference is computed as `mf` (see Section ??). This section specifies all hyperparameters of the OVOS calculation, ensuring reproducibility.

Equal for all OVOS calculations, the orbital optimization parameters are specified, such that `num_opt_virtual_orbs` specifies the number of active virtual spin-orbitals to optimize, `mo_coeff` sets the orbital coefficients, and `init_orbs` selects the initial orbital type (RHF or UHF).

Convergence is monitored via two criteria: a change in correlation energy $|\Delta E_{\text{corr}}| < 10^{-8}$ Hartree and a change in gradient norm $|\Delta G_{\text{corr}}| < 10^{-6}$ Hartree. Optimization proceeds for a maximum of 1000 iterations. Once convergence is achieved, consistent convergence is verified over `keep_track_max` = 50 consecutive iterations to ensure stability.

4.2 VQE Calculation Settings

All VQE calculations were performed using the SlowQuant branch of the Unrestricted VQE codebase [21]. The same molecular geometries and basis sets as described in Sections ?? and ?? were used for all VQE calculations, ensuring a direct comparison with OVOS results. The VQE calculations were executed on the same hardware environment as the OVOS calculations.

"Detail ansatz used, number of parameters, optimisation method, convergence criteria, etc."

"Give a flowchart of the VQE calculation steps, from Hamiltonian construction to energy evaluation and parameter update."

Chapter 5

Results

In this chapter, the results of the OVOS implementation are presented. The chapter is structured as follows:

- **Section 6.1** validates the implementation by comparing full-space MP2 correlation energies against PySCF and literature values.
- **Section 6.2** presents the main result: the fraction of MP2 correlation energy recovered as a function of the active virtual orbital count.
- **Section 6.3** analyses the convergence behaviour of the OVOS algorithm.
- **Section 6.4** compares different initialization strategies and their impact on final energies and convergence.
- **Section 6.5** investigates how the choice of basis set affects energy recovery.
- **Section 6.6** provides a direct comparison with the results published by Adamowicz and Bartlett in their original OVOS paper.
- **Section 6.7** assesses the computational cost and scaling of the OVOS algorithm.
- **Section 6.8** examines the T1 amplitudes in the converged OVOS orbitals to test for Brueckner-like properties.
- **Section 6.9** estimates the practical benefits of using OVOS orbitals as a reference state for quantum computing applications, including overlap with FCI states and resource estimates for VQE.

"Dont know if this list is too much..."

5.1 OVOS

...

5.1.1 Energy Recovery as a Function of Active Virtual Orbital Count

The results of the OVOS algorithm for all test molecules in the cc-pVDZ basis sets for start guess RHF, are presented in this section. Explore the correlation energy recovered by OVOS as a function of the number of active unoccupied orbitals $N'_{\text{act.unocc}}$.

CO

We observe that OVOS nearly steadily recovers increasingly more of the total correlation energy $E_{\text{corr}}^{\text{OVOS}}$ as $N'_{\text{act.unocc}}$ increases, with the percentage of correlation energy recovered reaching 100% at $N'_{\text{act.unocc}} = 21$. Reaching 90% recovery at $N'_{\text{act.unocc}} = 13$, which is approximately 62% of the total number of active unoccupied orbitals. The OVOS algorithm results for CO in the cc-pVDZ basis through each $N'_{\text{act.unocc}}$ to full space of active unoccupied orbitals are shown in Figure 5.1.

We mark the difference the OVOS algorithm make, amount of correlation energy recovered, at each $N'_{\text{act.unocc}}$ by an initial point with iterations to the final point for each $N'_{\text{act.unocc}}$ in Figure 5.1 as the points for each $N'_{\text{act.unocc}}$. The most recovered correlation energy is between $N'_{\text{act.unocc}} = 3$ and $N'_{\text{act.unocc}} = 10$ with $\Delta E_{\text{corr}}^{\text{OVOS}} \simeq 0.1$ Hartree. Convergence is nearly achieved in all cases, $N'_{\text{act.unocc}} = 6$ being non converged reaching maximum iterations, see Appendix ?? . Iteration counts show no clear pattern as $N'_{\text{act.unocc}}$ reaches the total number of active unoccupied orbitals. The correlation energy iterations are plottet in Appendix ?? to show the convergence behaviour of the OVOS algorithm for CO in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$.

Shortcomings of the OVOS algorithm are observed at $N'_{\text{act.unocc}} = 11, 12$ and $N'_{\text{act.unocc}} = 13$ for start guess RHF, where the algorithm fails to converge at an agreeable correlation energy, and instead converges to a local minimum with a much higher correlation energy. This suggest that the orbital optimization landscape has multiple local minima, which can trap the algorithm if the initial guess is not close enough to the global minimum. This is further explored in Section 6.4, where the effect of different initialisation strategies is investigated.

H2O

OVOS recovers increasingly more of the total correlation energy $E_{\text{corr}}^{\text{OVOS}}$ as $N'_{\text{act.unocc}}$ increases, with the percentage of correlation energy recovered reaching 100% at $N'_{\text{act.unocc}} = 19$. Reaching 90% recovery at $N'_{\text{act.unocc}} = 13$, which is approximately 68% of the total number of active unoccupied orbitals. The OVOS algorithm results for H_2O in the cc-pVDZ basis through each $N'_{\text{act.unocc}}$ to full space of active unoccupied orbitals are shown in Figure 5.2.

The diffrence OVOS makes for H_2O at each $N'_{\text{act.unocc}}$ is marked in Figure 5.2. The most recovered correlation energy is for $N'_{\text{act.unocc}} \in \{3, 4, 9, 10\}$ with $\Delta E_{\text{corr}}^{\text{OVOS}} \simeq 0.1$ Hartree. Convergence is achieved in nearly all cases, except for $N'_{\text{act.unocc}} = 13$, which is observed to stop at a correlation energy between $N'_{\text{act.unocc}} = 12$ and $N'_{\text{act.unocc}} = 14$, as seen in Figure 5.2. Iteration counts show no clear pattern as $N'_{\text{act.unocc}}$ increases, with the highest iteration count at $N'_{\text{act.unocc}} = 14$ for converged start guess RHF. The correlation energy iterations are plottet in Appendix ?? to show the convergence behaviour of the OVOS algorithm for H2O in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$.

HF

OVOS recovers increasingly more of the total correlation energy $E_{\text{corr}}^{\text{OVOS}}$ as $N'_{\text{act.unocc}}$ increases, with the percentage of correlation energy recovered reaching 100% at $N'_{\text{act.unocc}} = 14$. Reaching 90% recovery at $N'_{\text{act.unocc}} = 9$, which is approximately 64% of the total number of active unoccupied orbitals. The OVOS algorithm results for




Figure 5.1: Convergence of the MP2 correlation energy for CO in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$. The OVOS algorithm takes RHF as start guess. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated. "THIS PLOT NEEDED TO BE CORRECTED TO ONLY SHOW start guess RHF..."




Figure 5.2: Convergence of the MP2 correlation energy for H2O in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$. The OVOS algorithm takes RHF as start guess. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.




Figure 5.3: Convergence of the MP2 correlation energy for HF in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$. The OVOS algorithm takes RHF as start guess. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

HF in the cc-pVDZ basis through each $N'_{\text{act.unocc}}$ to full space of active unoccupied orbitals are shown in Figure 5.3.

The most correlation energy recovered by OVOS for HF is at $N'_{\text{act.unocc}} = 3, 4, 9$ with $\Delta E_{\text{corr}}^{\text{OVOS}} \simeq 0.625$ Hartree. Convergence is achieved in all cases, with iteration counts increasing and decreasing as $N'_{\text{act.unocc}}$ increases. The correlation energy iterations are plotted in Appendix ?? to show the convergence behaviour of the OVOS algorithm for HF in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$.

NH3

OVOS recovers increasingly more of the total correlation energy $E_{\text{corr}}^{\text{OVOS}}$ as $N'_{\text{act.unocc}}$ increases, except the convergence behavior for NH3 is less straightforward, we observe outliers at higher values of $N'_{\text{act.unocc}}$. It still holds that the percentage of correlation energy recovered reaches 100% at $N'_{\text{act.unocc}} = 24$, and we reach 90% recovery at $N'_{\text{act.unocc}} = 12$, which is approximately 50% of the total number of active unoccupied orbitals. The OVOS algorithm results for NH3 in the cc-pVDZ basis through each $N'_{\text{act.unocc}}$ to full space of active unoccupied orbitals are shown in Figure 5.4.

The most correlation energy recovered by OVOS for NH3 is at $N'_{\text{act.unocc}} = 3, \dots, 9$ with $\Delta E_{\text{corr}}^{\text{OVOS}} \simeq 0.75$ Hartree. Only at $N'_{\text{act.unocc}} = 8$ does OVOS not converge, which is observed to stop at an agreeable correlation energy, between $N'_{\text{act.unocc}} = 9$ and




Figure 5.4: Convergence of the MP2 correlation energy for NH3 in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$. The OVOS algorithm takes RHF as start guess. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

$N'_{\text{act.unocc}} = 10$, as seen in Figure 5.4. Iteration counts show no clear pattern as $N'_{\text{act.unocc}}$ increases, with the highest iteration count at $N'_{\text{act.unocc}} = 23$ of 797, undergoing both oscillations and divergence. The correlation energy iterations are plotted in Appendix ?? to show the convergence behaviour of the OVOS algorithm for NH3 in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$.

Li2

OVOS increasingly recovers more of the total correlation energy $E_{\text{corr}}^{\text{OVOS}}$ as $N'_{\text{act.unocc}}$ increases, quite slowly as $N'_{\text{act.unocc}}$ approaches 9, hereafter convergence issues show up, as multiple points deviate from the expected correlation energy behavior toward full space. The percentage of correlation energy recovered reaches 100% at $N'_{\text{act.unocc}} = 25$, and we reach 90% recovery at $N'_{\text{act.unocc}} = 7$, which is approximately 28% of the total number of active unoccupied orbitals. The OVOS algorithm results for Li2 in the cc-pVDZ basis through each $N'_{\text{act.unocc}}$ to full space of active unoccupied orbitals are shown in Figure 5.5.

The most correlation energy recovered by OVOS for Li2 is at $N'_{\text{act.unocc}} = 5, 6, 7$ with $\Delta E_{\text{corr}}^{\text{OVOS}} \approx 0.0125$ Hartree. OVOS converges for all $N'_{\text{act.unocc}}$. Iteration counts show no clear pattern as $N'_{\text{act.unocc}}$ increases, with the highest iteration count at $N'_{\text{act.unocc}} = 9$ of 569, undergoing both oscillations and divergence. The correlation energy iterations




Figure 5.5: Convergence of the MP2 correlation energy for Li2 in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$. The OVOS algorithm takes RHF as start guess. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

are plotted in Appendix ?? to show the convergence behaviour of the OVOS algorithm for Li2 in the cc-pVDZ basis at each $N'_{\text{act.unocc}}$.

Effect of Initialization Strategy -> Write it into above!!!

As the OVOS algorithm optimises the orbitals to recover correlation energy, the choice of initial orbitals can influence the convergence path and final energy. In this section, I compare three initialisation strategies: default SCF (RHF), warm-start from a previous N' run, and random unitary sampling. For the random sampling, I generate 1000 random unitary transformations of the initial orbitals and run OVOS for each, recording the final correlation energy. The best result from the random sampling is then compared to the default and warm-start strategies. The results for the benchmark molecules in the cc-pVDZ basis are presented below.

CO

The three initialisation strategies show a similar trend in the correlation energy recovered as $N'_{\text{act.unocc}}$ increases. It shows that while warm-starting from the previous N' run can provide an improvement in convergence where RHF initialisation struggles, itself can also struggle at certain $N'_{\text{act.unocc}}$, such as $N'_{\text{act.unocc}} = 12, \dots, 17$, where it con-




Figure 5.6: Comparison of final OVOS correlation energies for three initialisation strategies: default SCF (RHF), warm-start from previous N' orbitals, and best-of-1000 random unitary samples, for [molecule]/[basis]. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

verges to a local minimum with a much higher correlation energy. The best-of-random sampling strategy finds ...

All three methods deviate between each other for amount of correlation energy recovered at each $N'_{\text{act.unocc}}$, with the best-of-random sampling consistently finding at times significantly higher initial correlation energies than the other two strategies. Oppositely the warm-start strategy shows a smaller initial correlation energy than the default SCF (RHF) strategy at points the previous N' run converges to a local minimum.

Between the three strategies, we can pick the best result at each $N'_{\text{act.unocc}}$, which is shown as the green line in Figure 5.6. This best-of strategy shows a significant improvement in the correlation energy recovered at each $N'_{\text{act.unocc}}$, especially at points where the default and warm-start strategies struggle. This suggests that the OVOS algorithm has multiple local minima in the orbital rotation landscape, and that the choice of initial orbitals can significantly affect the final energy.

H2O

"Comment on the resulting data illustrated in Figure 5.7..."




Figure 5.7: Comparison of final OVOS correlation energies for three initialisation strategies: default SCF (RHF), warm-start from previous N' orbitals, and best-of-1000 random unitary samples, for [molecule]/[basis]. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.




Figure 5.8: Comparison of final OVOS correlation energies for three initialisation strategies: default SCF (RHF), warm-start from previous N' orbitals, and best-of-1000 random unitary samples, for [molecule]/[basis]. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

HF

"Comment on the resulting data illustrated in Figure 5.8..."

NH3

"Comment on the resulting data illustrated in Figure 5.9..."

5.1.2 Basis Set Dependence

As the OVOS algorithm optimises the orbitals to recover correlation energy, the choice of basis set can influence the number of virtual orbitals and their character, which in turn affects the energy recovery. In this section, I compare the energy recovery curves for different basis sets (6–31G, cc-pVDZ) for the benchmarked molecules. We use the best results of initialisation strategies based on the results from the previous section.




Figure 5.9: Comparison of final OVOS correlation energies for three initialisation strategies: default SCF (RHF), warm-start from previous N' orbitals, and best-of-1000 random unitary samples, for [molecule]/[basis]. Dashed red line marks the MP2 correlation energy at full space, PySCF reference. Dashed light blue line marks the OOMP2 correlation energy at full space, PySCF reference. Dashed purple line marks the FCI correlation energy at full space, PySCF reference. The points for each $N'_{\text{act.unocc}}$ show the change in correlation energy achieved by OVOS at each active virtual orbital count. A line is drawn to connect the points for visualization purposes, at which the number of iterations to convergence is annotated.

CO

Observing the correlated energy curves for CO in the 6–31G (11) and cc–pVDZ (21) basis sets, we see that the larger cc–pVDZ basis set, which has more virtual orbitals, allows for a higher fraction of correlation energy to be recovered at lower $N'_{\text{act.unocc}}$. At $N'_{\text{act.unocc}} = 25\%$ of the total active virtual orbitals, the cc–pVDZ basis recovers less than 6–31G basis, for which as $N'_{\text{act.unocc}}$ increases, the two basis sets show similar trends. The energy recovery curves for CO in the 6–31G and cc–pVDZ basis sets are shown in Figure 5.10.

If instead we looked at an fixed number of active virtual orbitals, $N'_{\text{act.unocc}} = 5$, their percentages of the total active virtual orbitals are different for the two basis sets, with 6–31G having a higher percentage than cc–pVDZ. We can find this point by counting 5 points along the line for each basis set in Figure 5.10. At this fixed number of active virtual orbitals, the 6–31G basis recovers a higher fraction of correlation energy than the cc–pVDZ basis, as the total number of virtual orbitals is smaller in 6–31G.




Figure 5.10: Basis set dependence of energy recovery for CO. The MP2 energy correlation is plotted against the active virtual orbitals for the basis sets: 6–31G (stripped) and cc–pVDZ (full) line.

H2O

The correlated energy curves for H2O in the 6–31G (8) and cc–pVDZ (19) basis sets show a stronger trend in the larger cc–pVDZ basis set, beating the smaller 6–31G basis set at most points along the curve, see Figure 5.11. At $N'_{\text{act.unocc}} = 25\%$ of the total active virtual orbitals, the cc–pVDZ basis recovers more correlation energy than the 6–31G basis. As $N'_{\text{act.unocc}}$ increases, the two basis sets show similar trends, with the cc–pVDZ basis meeting the 6–31G basis. At a fixed number of active virtual

orbitals, $N'_{\text{act.unocc}} = 5$, the 6-31G basis recovers a higher correlation energy than the cc-pVDZ basis, as the total number of virtual orbitals is smaller in 6-31G.




Figure 5.11: Basis set dependence of energy recovery for H2O. The MP2 energy correlation is plotted against the active virtual orbitals for the basis sets: 6-31G (stripped) and cc-pVDZ (full) line.

HF

The correlated energy curves for HF in the 6-31G (6) and cc-pVDZ (14) basis sets show a stronger trend in the larger cc-pVDZ basis set, beating the smaller 6-31G basis set at most points along the curve, see Figure 5.12. At $N'_{\text{act.unocc}} = 25\%$ of the total active virtual orbitals, the cc-pVDZ basis recovers more correlation energy than the 6-31G basis. As $N'_{\text{act.unocc}}$ increases, the two basis sets show similar trends, with the cc-pVDZ basis meeting the 6-31G basis. At a fixed number of active virtual orbitals, $N'_{\text{act.unocc}} = 5$, the 6-31G basis recovers a higher correlation energy than the cc-pVDZ basis, as the total number of virtual orbitals is smaller in 6-31G.

NH3

The correlated energy curves for NH3 in the 6-31G (10) and cc-pVDZ (24) basis sets show a stronger trend in the larger cc-pVDZ basis set, beating the smaller 6-31G basis set at most points along the curve, see Figure 5.13. At $N'_{\text{act.unocc}} = 25\%$ of the total active virtual orbitals, the cc-pVDZ basis recovers more correlation energy than the 6-31G basis. As $N'_{\text{act.unocc}}$ increases, the two basis sets show similar trends, with the cc-pVDZ basis meeting the 6-31G basis. At a fixed number of active virtual orbitals, $N'_{\text{act.unocc}} = 5$, the 6-31G basis recovers a higher correlation energy than the cc-pVDZ basis, as the total number of virtual orbitals is smaller in 6-31G.




Figure 5.12: Basis set dependence of energy recovery for HF. The MP2 energy correlation is plotted against the active virtual orbitals for the basis sets: 6–31G (stripped) and cc–pVDZ (full) line.

Li2

The correlated energy curves for Li2 in the 6–31G (15) and cc–pVDZ (25) basis sets show a similar trend, with the cc–pVDZ basis set generally recovering more correlation energy than the 6–31G basis set. This is consistent with the observations for previous molecules, where larger basis sets tend to perform better in terms of energy recovery, especially at lower $N'_{\text{act.unocc}}$. At $N'_{\text{act.unocc}} = 25\%$ of the total active virtual orbitals, the cc–pVDZ basis recovers more correlation energy than the 6–31G basis. As $N'_{\text{act.unocc}}$ increases, the two basis sets show similar trends, with the cc–pVDZ basis meeting the 6–31G basis. Where when they get to the bad points of cc–pVDZ, the 6–31G basis performs better. At a fixed number of active virtual orbitals, $N'_{\text{act.unocc}} = 5$, the 6–31G basis recovers a higher correlation energy than the cc–pVDZ basis, as the total number of virtual orbitals is smaller in 6–31G.

...

Comparison with Adamowicz and Bartlett (1987) - "Skip?"

Adamowicz & Bartlett's 1987 paper presents OVOS results for the CH₂ molecule in a specific basis set ... in their Table I ... I include their results here to show potential of OVOS at bigger basis sets, and to compare with the results obtained in this work. However, it is important to note that there are differences in the computational details between their work and this one, such as their handling of convergence (which does not seem to be detailed, e.g. ill-conditioned Hessian, iteration count, ...), the version of the basis set used, and the geometry of the molecule. These differences can lead to variations in the correlation energy values, making a direct comparison challenging.




Figure 5.13: Basis set dependence of energy recovery for NH₃. The MP2 energy correlation is plotted against the active virtual orbitals for the basis sets: 6–31G (stripped) and cc–pVDZ (full) line.

Looking past not being able to replicate the exact setup of Adamowicz & Bartlett, the basis sets behaviour shows some similarities with their findings, the energy recovery curve showing a rapid increase in correlation energy recovered as $N'_{\text{act.unocc}}$ increases 5.1.1, the increase in basis set size leading to a higher correlation energy recovery at the same $N'_{\text{act.unocc}}$.

5.2 VQE

...

5.2.1 OVOS Orbitals as a VQE Reference State

"Comment on VQE results for ... in the ... basis, with the OVOS orbitals as reference state, compared to the UHF orbitals, and UMP2 natural orbitals as reference state. Comment on iteration and energy in VQE calculations using orbital optimization or not. Comment on random initialization of thetas vs previous geometry final values... Comment on the qubit savings and parameter savings... "

5.2.2 T1 Amplitudes and the Brueckner Orbital Connection

In this section, we investigate the connection between the OVOS-optimised orbitals and Brueckner orbitals, by the maximum-overlap criterion and Brillouin-Brueckner condition. Brueckner orbitals are defined as the orbitals for which the single excitation amplitudes (T1) vanish, i.e. the excited determinants with one electron promoted from




Figure 5.14: Basis set dependence of energy recovery for Li₂. The MP2 energy correlation is plotted against the active virtual orbitals for the basis sets: 6–31G (stripped) and cc–pVDZ (full) line.

an occupied to a virtual orbital do not contribute to the wavefunction, see Section ?? . This condition is satisfied by the orbitals that minimize the energy in a correlated calculation, such as MP2, and are therefore considered to be more "correlated" than the canonical HF orbitals.

I report here having observed computed T1 amplitudes for the OVOS-optimised orbitals that are significantly small, on the order of 10^{-16} , for all the test molecules and basis sets considered, which is consistent with the Brueckner condition. In contrast, the T1 amplitudes for the canonical HF orbitals are typically ...

"Look at overlap of FCI wavefunction with the reference state, for both HF and OVOS orbitals, as a function of the number of active virtual orbitals, to see if the OVOS orbitals have a higher overlap with the FCI wavefunction than the HF orbitals, which would be another indication that they are closer to Brueckner orbitals. Refer to plot of orbitals as a heatmap in the appendix."

5.3 ooVQE

...

Chapter 6

Conclusion and Outlook

...

6.1 Conclusion

...

6.1.1 Summary of Findings

...

6.1.2 Critical Assessment

...

6.1.3 Implications for Quantum Computing

...

6.2 Outlook

...

6.2.1 Algorithmic Improvements

HESSIAN

NEWTON STEP VS DIIS

6.2.2 Method Extensions

AS Bartlett ALSO COVER TO HIGHER ORDER...

6.2.3 Quantum Computing Applications

OVOS Orbitals as a VQE Reference State

ON QUANTUM HARDWARE!!!!!!

Brueckner Orbitals and Circuit Depth on NISQ Devices

Qubit Count Reduction and Fault-Tolerant Perspectives

6.2.4 Larger Molecular Systems

EXPAND TO LARGER BASIS SETS OR BIGGER MOLECULES - PROTEINS HOW DOES OVOS PERFORM...

6.2.5 Benchmarking and Validation

EXPAND THE BENCHMARKING AGAINST OTHER METHODS ...

Bibliography

- [1] C. David Sherrill and Henry F. Schaefer. The configuration interaction method: Advances in highly correlated approaches. In Per-Olov Löwdin, editor, *Advances in Quantum Chemistry*, volume 34, pages 143–269. Academic Press, 1999. doi: 10.1016/S0065-3276(08)60532-8.
- [2] Chr. Møller and M. S. Plesset. Note on an approximation treatment for many-electron systems. *Phys. Rev.*, 46:618–622, 1934. doi: 10.1103/PhysRev.46.618.
- [3] Trygve Helgaker, Poul Jørgensen, and Jeppe Olsen. *Molecular Electronic-Structure Theory*. Wiley, 2000. doi: 10.1002/9781119019572.
- [4] Alberto Peruzzo, Jarrod McClean, Peter Shadbolt, Man-Hong Yung, Xiao-Qi Zhou, Peter J. Love, Alán Aspuru-Guzik, and Jeremy L. O’Brien. A variational eigenvalue solver on a photonic quantum processor. *Nat. Commun.*, 5:4213, 2014. doi: 10.1038/ncomms5213.
- [5] Yudong Cao, Jonathan Romero, Jonathan P. Olson, Matthias Degroote, Peter D. Johnson, Maria Kieferova, Ian D. Kivlichan, Tim Menke, Borja Peropadre, Nicolas P. D. Sawaya, Sukin Sim, Libor Veis, and Alán Aspuru-Guzik. Quantum chemistry in the age of quantum computing. *Chem. Rev.*, 119:10856–10915, 2019. doi: 10.1021/acs.chemrev.8b00803.
- [6] L. Adamowicz and R. J. Bartlett. Optimized virtual orbital space for high-level correlated calculations. *J. Chem. Phys.*, 86(11):6314–6324, 1987. doi: 10.1063/1.452468.
- [7] Per-Olov Löwdin. Quantum theory of many-particle systems. I. physical interpretations by means of density matrices, natural spin-orbitals, and convergence problems in the method of configurational interaction. *Phys. Rev.*, 97:1474–1489, 1955. doi: 10.1103/PhysRev.97.1474.
- [8] Andrew G. Taube and Rodney J. Bartlett. Frozen natural orbitals: Systematic basis set truncation for coupled-cluster theory. *Collect. Czech. Chem. Commun.*, 70:837–850, 2005. doi: 10.1135/cccc20050837.
- [9] Qiming Sun, Xing Zhang, Samragni Banerjee, Peng Bao, Marc Barbry, Nick S. Blunt, Nikolay A. Bogdanov, George H. Booth, Jia Chen, Zhi-Hao Cui, Janus Juul Eriksen, Yang Gao, Sheng Guo, Jan Hermann, Matthew R. Hermes, Kevin Koh, Peter Koval, Susi Lehtola, Zhendong Li, Junzi Liu, Narbe Mardirossian, James D. McClain, Mario Motta, Bastien Mussard, Hong-Zhou Pham, Artem Pulkin, Wirawan Purwanto, Paul J. Robinson, Enrico Ronca, Elvira R. Sayfutyarova, Maximilian Scheurer, Henry F. Schurkus, James E. T. Smith, Chong Sun, Shi-Ning Sun, Shiv Upadhyay, Lucas K. Wagner, Xiao Wang, Alec White, James Daniel Whitfield, Mark J. Williamson, Sebastian Wouters, Jun Yang, Jason M. Yu, Tianyu Zhu, Timothy C. Berkelbach, Sandeep Sharma, Alexander Yu. Sokolov, and Garnet Kin-Lic Chan. Recent developments in the PySCF program package. *J. Chem. Phys.*, 153:024109, 2020. doi: 10.1063/5.0006074.

- [10] Jacob T. Seeley, Martin J. Richard, and Peter J. Love. The Bravyi–Kitaev transformation for quantum computation of electronic structure. *J. Chem. Phys.*, 137:224109, 2012. doi: 10.1063/1.4768229.
- [11] Joonho Lee and Martin Head-Gordon. Regularized orbital-optimized second-order Møller–Plesset perturbation theory: A reliable fifth-order scaling electron correlation model with orbital energy dependent regularizers. *J. Chem. Theory Comput.*, 14:5203–5219, 2018. doi: 10.1021/acs.jctc.8b00731. arXiv:1807.06185.
- [12] Attila Szabo and Neil S. Ostlund. *Modern Quantum Chemistry: Introduction to Advanced Electronic Structure Theory*. Dover Publications, 1989. ISBN 978-0-486-69186-2.
- [13] Frank Jensen. *Introduction to Computational Chemistry*. Wiley, 3rd edition, 2017. ISBN 978-1-118-82599-0.
- [14] R. Ditchfield, W. J. Hehre, and J. A. Pople. Self-consistent molecular-orbital methods. ix. an extended gaussian-type basis for molecular-orbital studies of organic molecules. *J. Chem. Phys.*, 54:724–728, 1971. doi: 10.1063/1.1674902.
- [15] W. J. Hehre, R. Ditchfield, and J. A. Pople. Self-consistent molecular orbital methods. xii. further extensions of gaussian-type basis sets for use in molecular orbital studies of organic molecules. *J. Chem. Phys.*, 56:2257–2261, 1972. doi: 10.1063/1.1677527.
- [16] James D. Dill and John A. Pople. Self-consistent molecular orbital methods. xv. extended gaussian-type basis sets for lithium, beryllium, and boron. *J. Chem. Phys.*, 62:2921–2923, 1975. doi: 10.1063/1.430801.
- [17] Thom H. Dunning. Gaussian basis sets for use in correlated molecular calculations. I. the atoms boron through neon and hydrogen. *J. Chem. Phys.*, 90:1007–1023, 1989. doi: 10.1063/1.456153.
- [18] Nicholas C. Handy, John A. Pople, Martin Head-Gordon, Krishnan Raghavachari, and Gary W. Trucks. Size-consistent Brueckner theory limited to single and double substitutions. *Chem. Phys. Lett.*, 164:185–192, 1989. doi: 10.1016/0009-2614(89)85013-4.
- [19] Keith A. Brueckner. Two-body forces and nuclear saturation. III. details of the structure of the nucleus. *Phys. Rev.*, 97:1353–1366, 1954. doi: 10.1103/PhysRev.97.1353.
- [20] George D. Purvis and Rodney J. Bartlett. A full coupled-cluster singles and doubles model: The inclusion of disconnected triples. *J. Chem. Phys.*, 76:1910–1918, 1982. doi: 10.1063/1.443164.
- [21] E. Kjellgren, K. M. Ziems, and P. Volsgaard. Slowquant unrestricted vqe, 2026.
- [22] Benjamin P. Pritchard, Doaa Altarawy, Brett Didier, Tara D. Gibsom, and Theresa L. Windus. A new basis set exchange: An open, up-to-date resource for the molecular sciences community. *J. Chem. Inf. Model.*, 59:4814–4820, 2019. doi: 10.1021/acs.jcim.9b00725.

Appendix A

Implementation Validation

lead into appendix A ...

A.1 Validation of the Implementation

PySCF gives me an opportunity to validate my implementation of full space MP2 by comparing with their built-in MP2 module, `pyscf.mp.MP2`. I present the deviation $\Delta E_{\text{corr}} = E_{\text{corr}}^{\text{code}} - E_{\text{corr}}^{\text{PySCF}}$ is within numerical precision for all test molecules in cc-pVDZ, see Figure A.1.

Molecule	$E_{\text{corr}}^{\text{OVOS}}$ (Hartree)	$E_{\text{corr}}^{\text{PySCF}}$ (Hartree)	ΔE_{corr} (Hartree)
CO	-2.90778E-01	-2.90778E-01	4.25617E-8
H ₂ O	-2.04003E-01	-2.04003E-01	2.29192E-9
HF	-2.03781E-01	-2.03781E-01	1.23463E-10
NH ₃	-1.88825E-01	-1.88825E-01	8.31300E-10
Li ₂	-2.00076E-02	-2.00076E-02	1.11200E-9

Table A.1: Validation of the OVOS implementation against PySCF mp.MP2. Deviations ΔE are within numerical precision ($|\delta E| < \approx 10^{-8}$ Hartree). All molecules are of the cc-pVDZ basis for this validation.

It is observed that CO's deviation is barely at acceptable range of 10^{-8} Hartree, on the other hand, the other molecules show deviations below the acceptable range, which confirms that my implementation of the MP2 correlation energy is correct, and that the OVOS code correctly reproduces the full-space MP2 limit when $N'_{\text{act.unocc}} = N_{\text{act.unocc}}$.

Verifying that the full-space MP2 energies match PySCF's results, any differences observed in the OVOS results are due to the orbital optimization and active space truncation, rather than a bug in the MP2 implementation.

Appendix B

Additional Data and Figures

B.1 Full MP2 Reference Energies

"I was thinking to do this in paint to make them sit together first!!!"

B.2 Complete OVOS Energy Recovery Tables

"I was thinking to do this in paint to make them sit together first!!!"

B.3 Convergence Plots for All Systems

"I was thinking to do this in paint to make them sit together first!!!"

Appendix C

Implementation Details and Source Code

C.1 Repository and Version Information

C.2 Spin-Orbital Integral Conversion

Listing C.1: Spatial-to-spin ERI conversion (`spatial_2_spin_eri_optimized`).

```
1 # Convert spatial ERIs to spin-orbital basis (interleaved alpha  
  /beta)  
2 def spatial_2_spin_eri_optimized(self, eri_aaaa, eri_aabb,  
  eri_bbbb):  
3     n_spin = self.tot_num_spin_orbs  
4     eri_spin = np.zeros((n_spin, n_spin, n_spin, n_spin))  
5     # alpha-alpha block (even indices)  
6     eri_spin[0::2, 0::2, 0::2, 0::2] = eri_aaaa  
7     # beta-beta block (odd indices)  
8     eri_spin[1::2, 1::2, 1::2, 1::2] = eri_bbbb  
9     # alpha-alpha / beta-beta cross block  
10    eri_spin[0::2, 0::2, 1::2, 1::2] = eri_aabb  
11    eri_spin[1::2, 1::2, 0::2, 0::2] = eri_aabb.transpose(2, 3,  
      0, 1)  
12    return eri_spin
```

C.3 MP2 Energy and Amplitude Computation

C.4 Orbital Optimisation: Gradient and Hessian

C.5 Profiling Results